

# **Real-time Financial Fraud Detection with Modern Python**

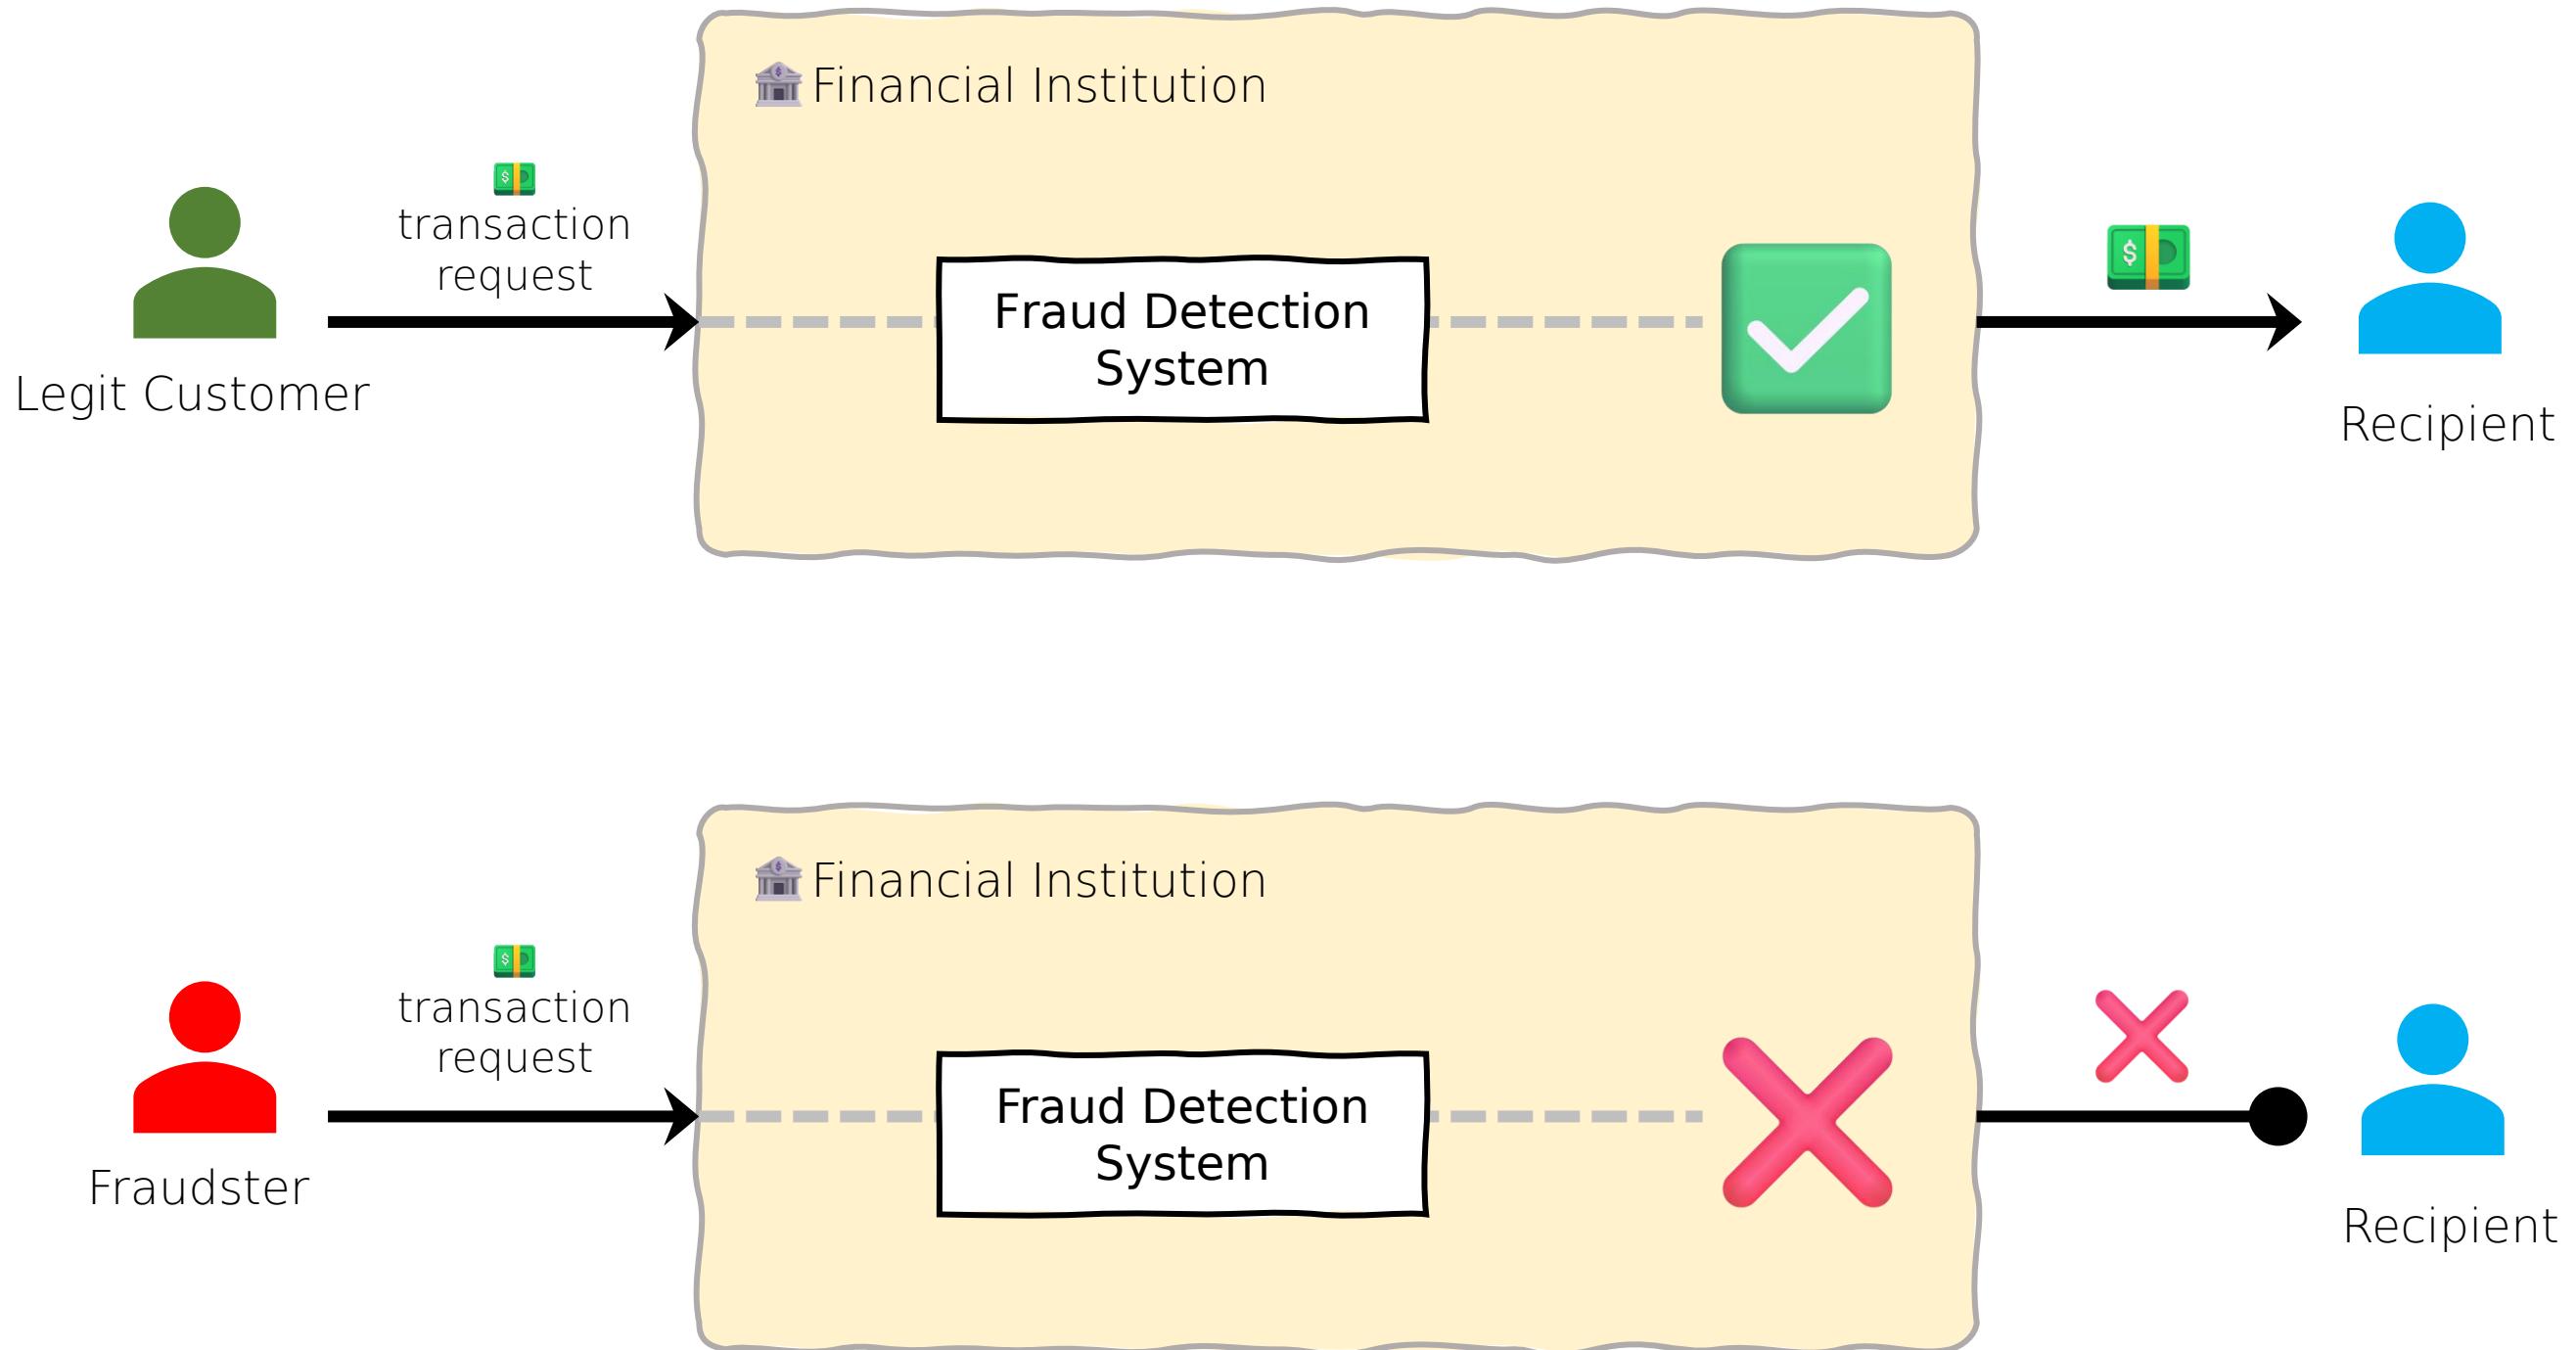
How to build production-grade fraud detection systems that make sub-second assessments on financial transactions while adversaries constantly adapt.

César Soto Valero  
Data Scientist / ML Engineer

# In This Session

1. Problem framing
2. Decision-first metrics
3. Time-aware validation
4. Latency-aware modeling
5. Shipping & operations

# **1. Problem Framing**





# The Stakes Are Very High!



## Real Money, Realtime

Every transaction requires a decision in milliseconds. Delays mean lost revenue or unhappy customers.



## Fraud Evolves Daily

Rule-based systems can't keep the pace with adversaries who iterate and adapt like well-funded startups.

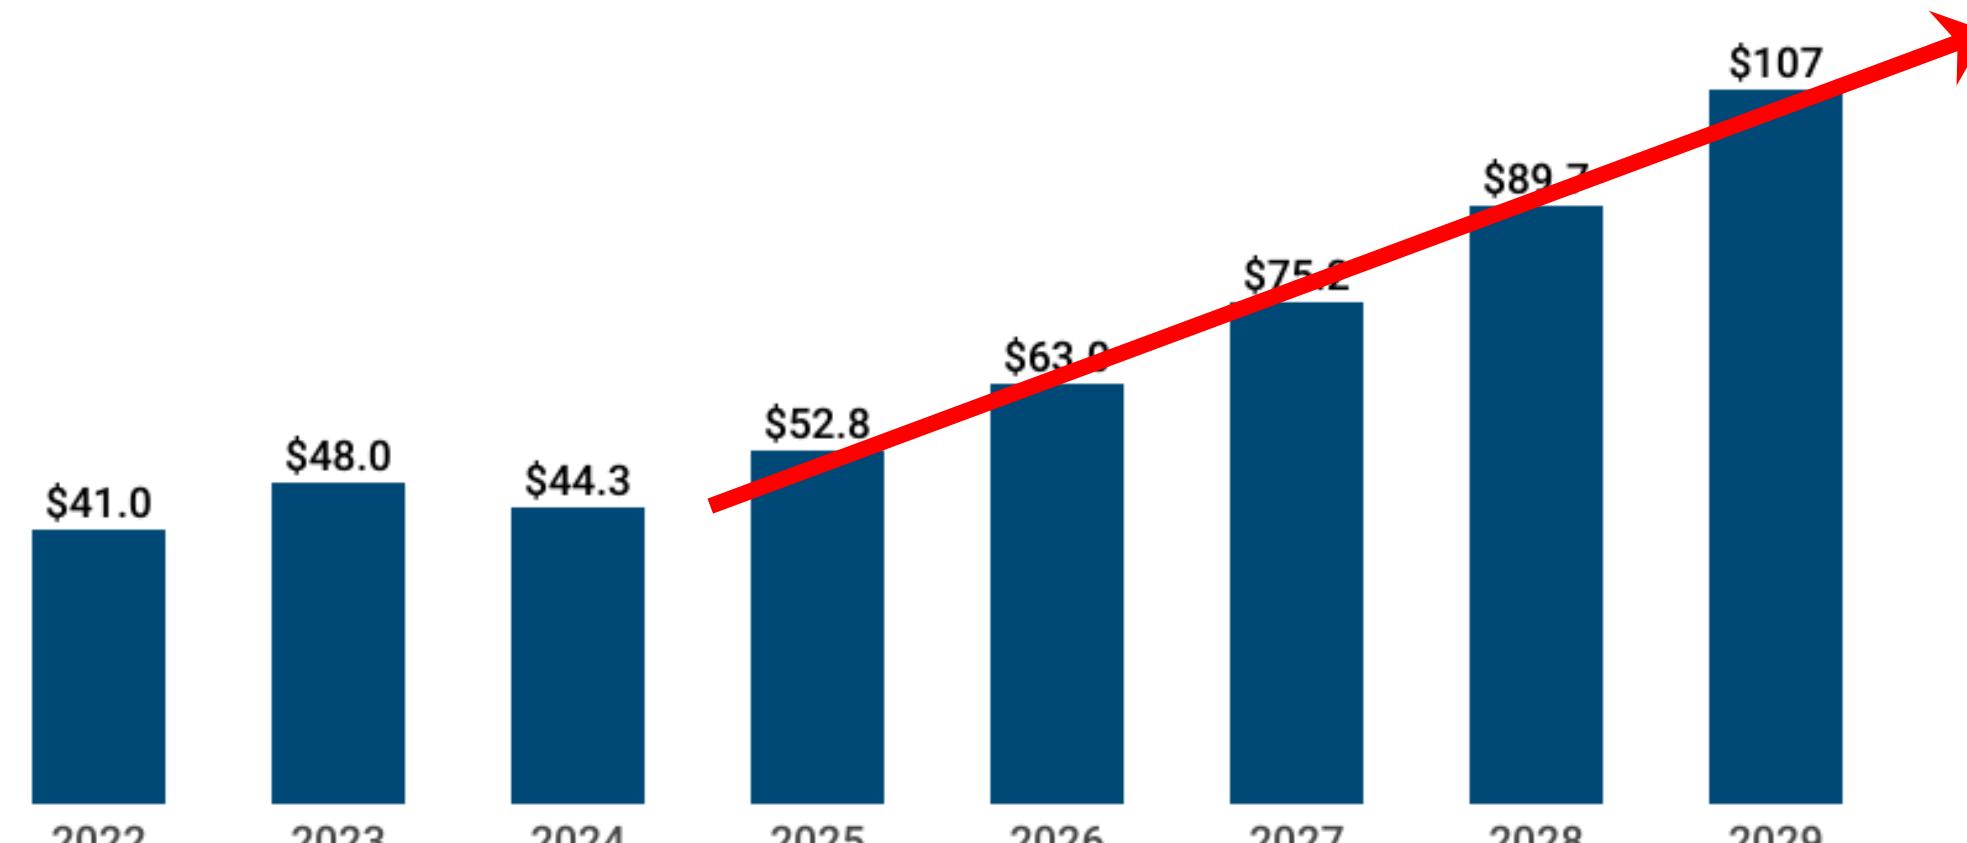


## Trust Is Fragile

False positives block legitimate users, eroding trust faster than fraud costs money.

# Global Losses to Online Payment Fraud

(in billions)



Source: Statista and Merchant Risk Council

Years 2025 onward are projections

# Why This Is Hard



## Extreme Class Imbalance

Fraud represents 0.1–2% of transactions in most systems. Finding needles in haystacks is the core technical challenge. Traditional accuracy metrics become meaningless.



## Delayed Ground Truth

Labels arrive days or weeks after decisions. Chargebacks and investigations take time. Training on outdated patterns while fraud tactics shift creates constant model drift.



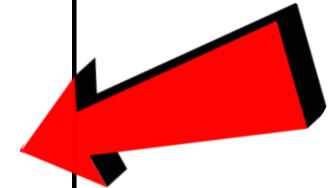
## Active Adversaries

Fraudsters deliberately probe your thresholds and test detection boundaries. They reverse-engineer your models through systematic experimentation, forcing constant adaptation.

## **2. Decision-First Metrics**

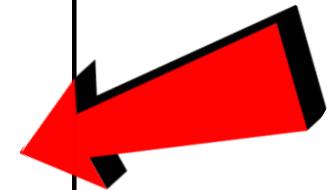
## GROUND TRUTH

		POSITIVE (fraudulent txt)	NEGATIVE (genuine txt)
		True Positive (TP)	False Positive (FP)
PREDICTION	POSITIVE (fraudulent txt)	<p><u>Reality:</u> Fraud <u>Prediction:</u> Fraud <u>Result:</u> Good</p>	<p><u>Reality:</u> Not Fraud <u>Prediction:</u> Fraud <u>Result:</u> Unnecessary costs and operational inefficiency</p>
	NEGATIVE (genuine txt)	<p><u>Reality:</u> Fraud <u>Prediction:</u> Not Fraud <u>Result:</u> Non-compliance and induced costs for the bank and society</p>	<p><u>Reality:</u> Not Fraud <u>Prediction:</u> Not Fraud <u>Result:</u> Good</p>



## GROUND TRUTH

		POSITIVE (fraudulent txt)	NEGATIVE (genuine txt)
		True Positive (TP)	False Positive (FP)
PREDICTION	POSITIVE (fraudulent txt)	<p><u>Reality:</u> Fraud <u>Prediction:</u> Fraud <u>Result:</u> Good</p>	<p><u>Reality:</u> Not Fraud <u>Prediction:</u> Fraud <u>Result:</u> Unnecessary costs and operational inefficiency</p>
	NEGATIVE (genuine txt)	<p><u>Reality:</u> Fraud <u>Prediction:</u> Not Fraud <u>Result:</u> Non-compliance and induced costs for the bank and society</p>	<p><u>Reality:</u> Not Fraud <u>Prediction:</u> Not Fraud <u>Result:</u> Good</p>



## GROUND TRUTH

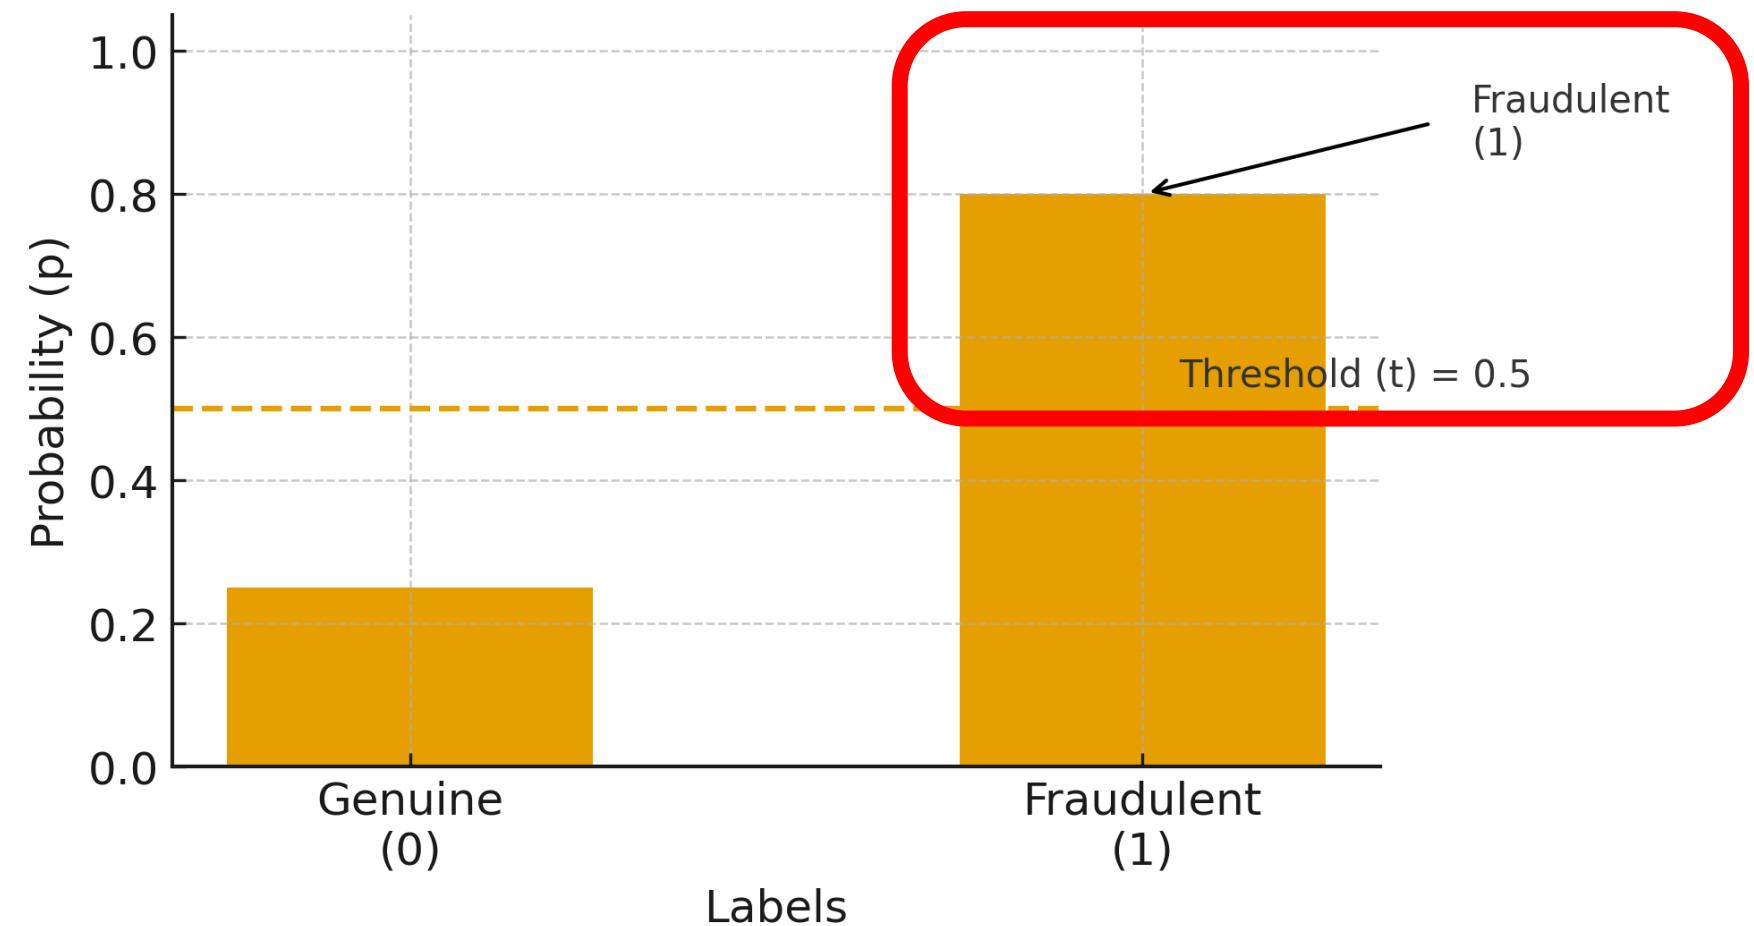
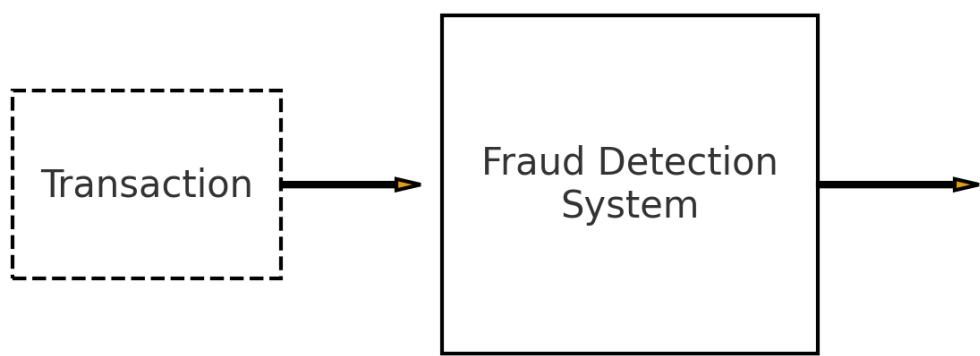
		POSITIVE (fraudulent txt)	NEGATIVE (genuine txt)
		True Positive (TP)	False Positive (FP)
PREDICTION	POSITIVE (fraudulent txt)	<p><u>Reality:</u> Fraud <u>Prediction:</u> Fraud <u>Result:</u> Good</p>	<p><u>Reality:</u> Not Fraud <u>Prediction:</u> Fraud <u>Result:</u> Unnecessary costs and operational inefficiency</p>
	NEGATIVE (genuine txt)	<p><u>Reality:</u> Fraud <u>Prediction:</u> Not Fraud <u>Result:</u> Non-compliance and induced costs for the bank and society</p>	<p><u>Reality:</u> Not Fraud <u>Prediction:</u> Not Fraud <u>Result:</u> Good</p>

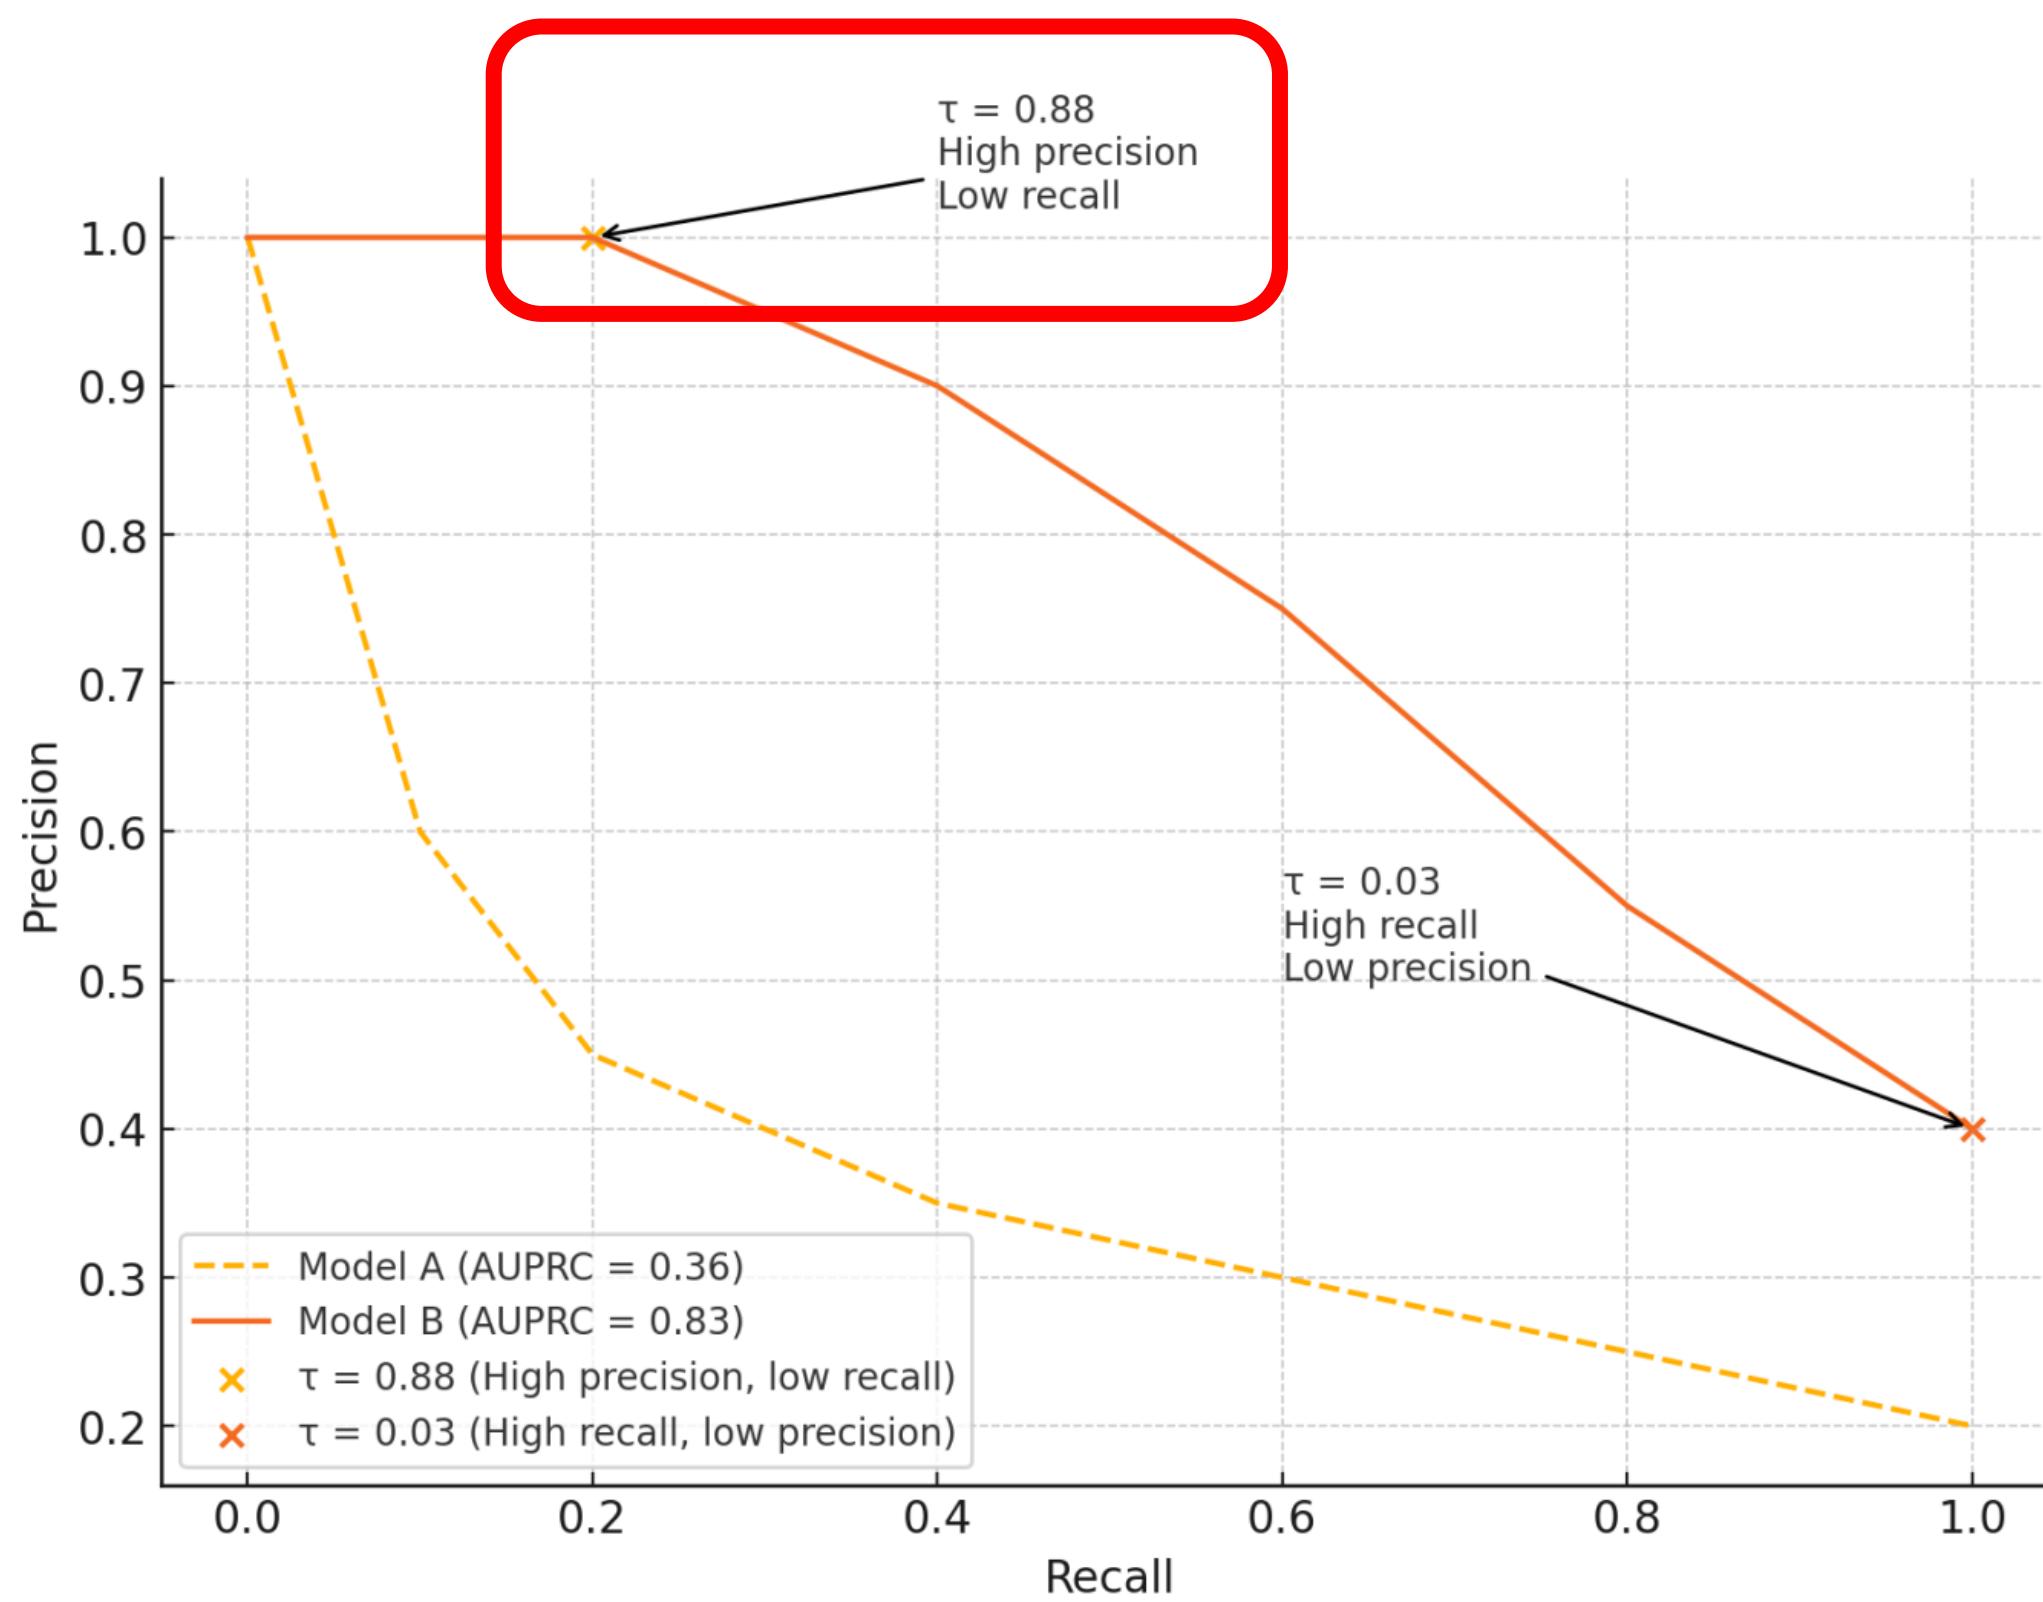


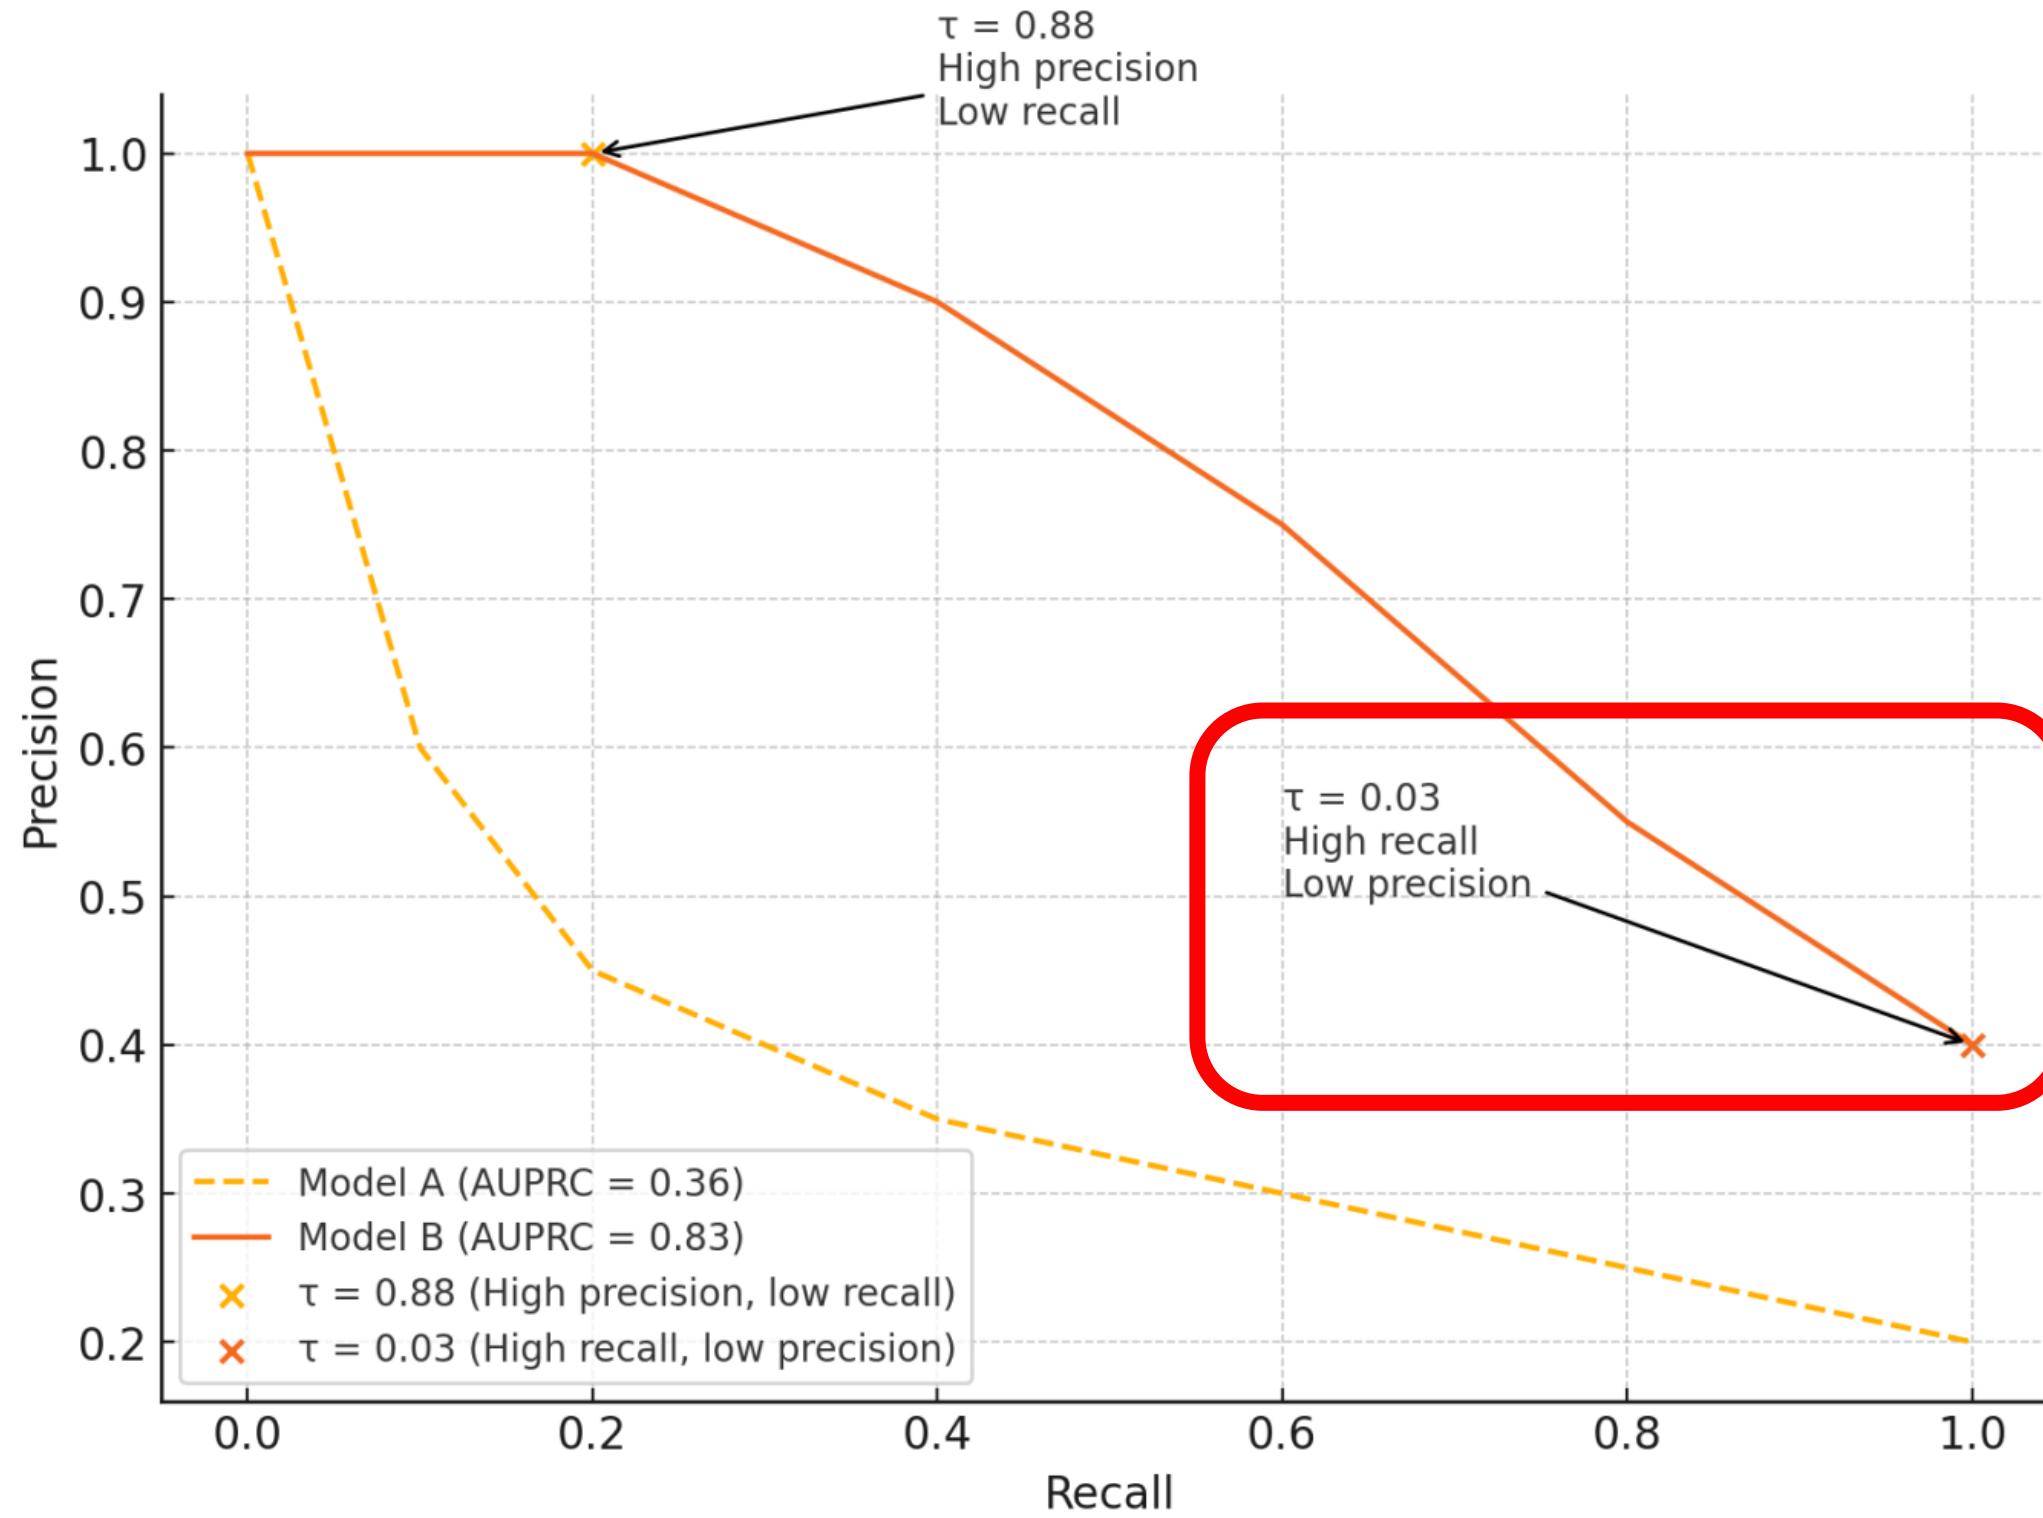
$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$F1 = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$









```
from sklearn.metrics import precision_score, recall_score, precision_recall_curve

# y_true: 0 = genuine, 1 = fraud
# y_pred: model's predicted labels after thresholding probabilities
precision = precision_score(y_true, y_pred, pos_label=1)
recall    = recall_score(y_true, y_pred, pos_label=1)

precision, recall, thresholds = precision_recall_curve(y_true, p_fraud, pos_label=1)
```

# Extreme Class Imbalance

A model that labels everything as legitimate achieves 99.9% accuracy while catching zero fraud. Accuracy is a dangerous illusion when classes are imbalanced.

## Precision and Recall Tell the Real Story

Precision measures how many alerts are actually fraud. Recall captures what percentage of fraud you catch. Both matter intensely for business outcomes.

## Thresholds Drive Business Impact

Moving a decision threshold changes precision-recall balance and directly affects revenue, review costs, and customer experience. The threshold is a business lever, not just a technical parameter.

### **3. Time-Aware Validation**

# Delayed Ground Truth

---

# Delayed Ground Truth

## Temporal Train/Val/Test Splits

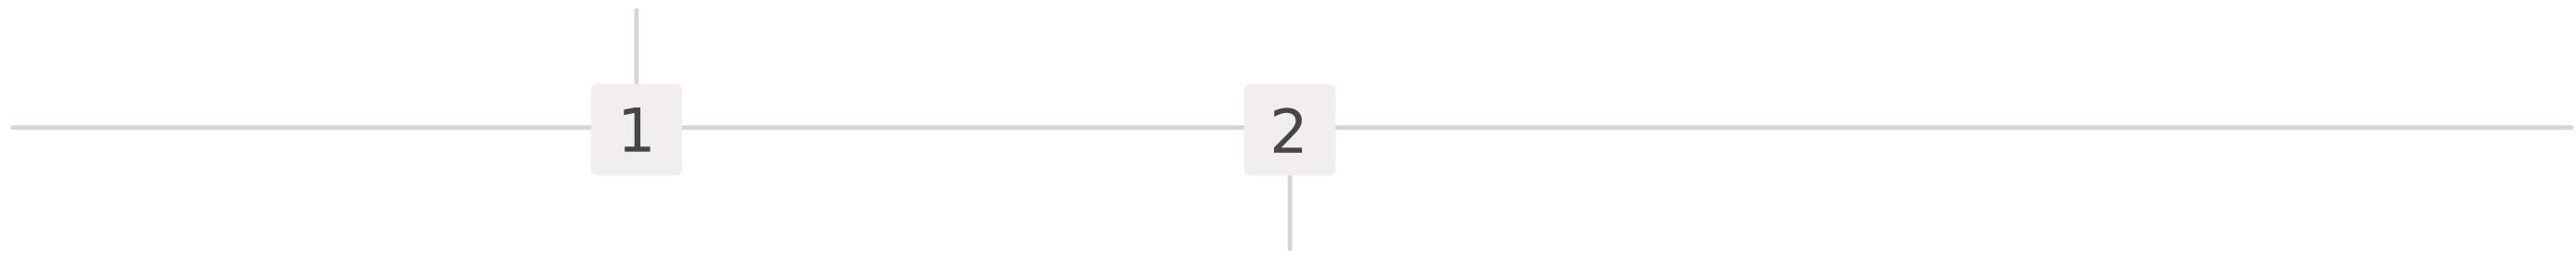
Never randomize timestamps. Train on past data, validate on near-future, test on held-out future. Respect time's arrow.



# Delayed Ground Truth

## Temporal Train/Val/Test Splits

Never randomize timestamps. Train on past data, validate on near-future, test on held-out future. Respect time's arrow.



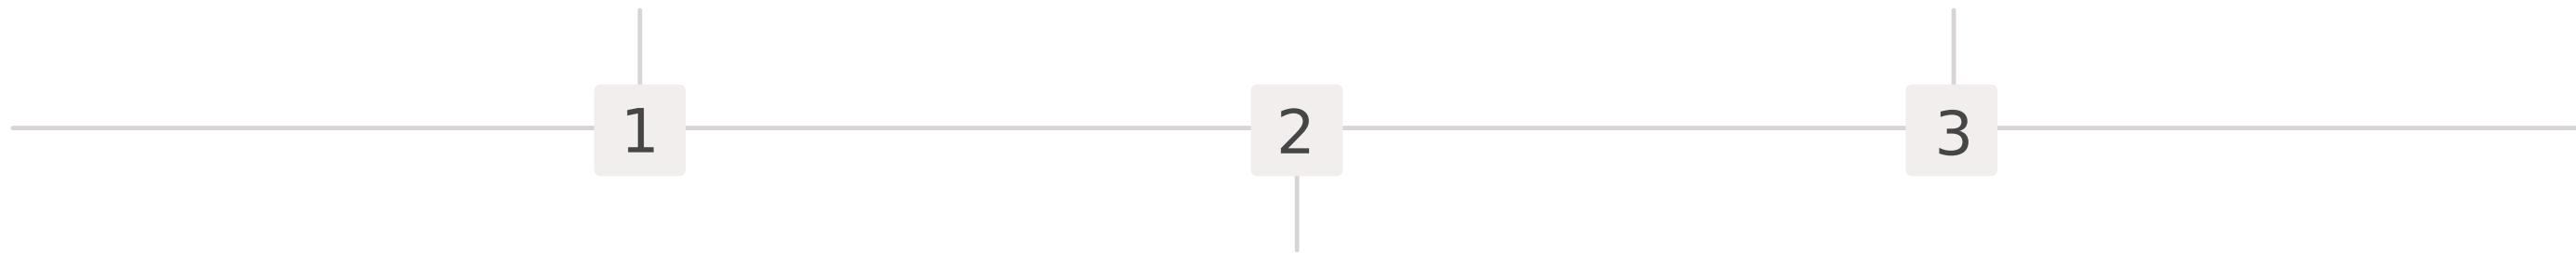
## Rolling Windows with Safety Gaps

Use sliding time windows that mimic production deployment. Add gaps between train and validation to simulate label delay.

# Delayed Ground Truth

## Temporal Train/Val/Test Splits

Never randomize timestamps. Train on past data, validate on near-future, test on held-out future. Respect time's arrow.



## Strictly Prevent Leakage Across Time

Future information must never leak into training. One timestamp error invalidates all results and creates false confidence.

## Rolling Windows with Safety Gaps

Use sliding time windows that mimic production deployment. Add gaps between train and validation to simulate label delay.

# Delayed Ground Truth

## Temporal Train/Val/Test Splits

Never randomize timestamps. Train on past data, validate on near-future, test on held-out future. Respect time's arrow.



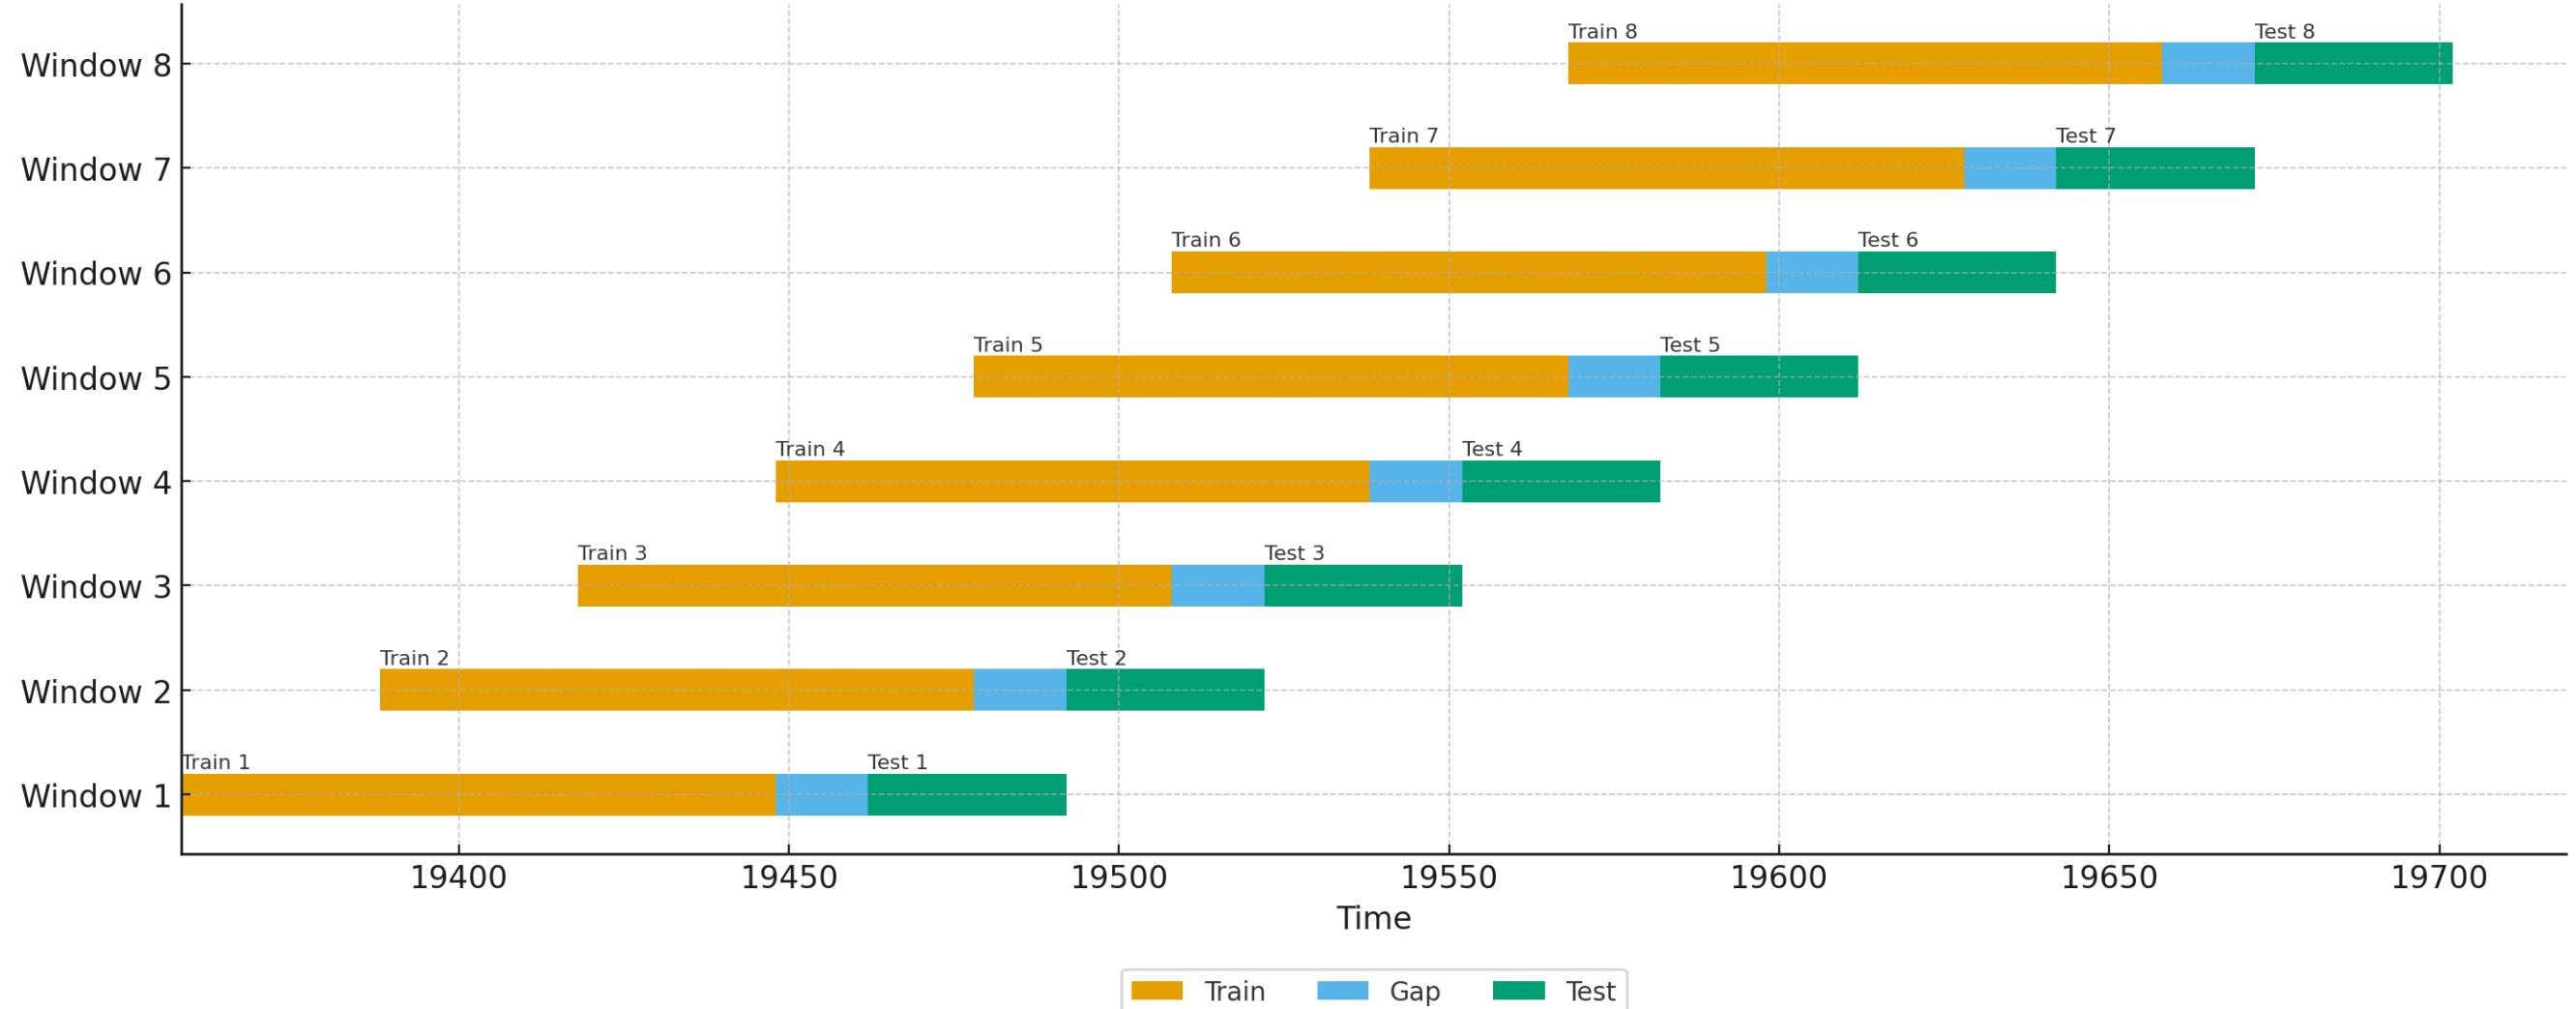
## Rolling Windows with Safety Gaps

Use sliding time windows that mimic production deployment. Add gaps between train and validation to simulate label delay.



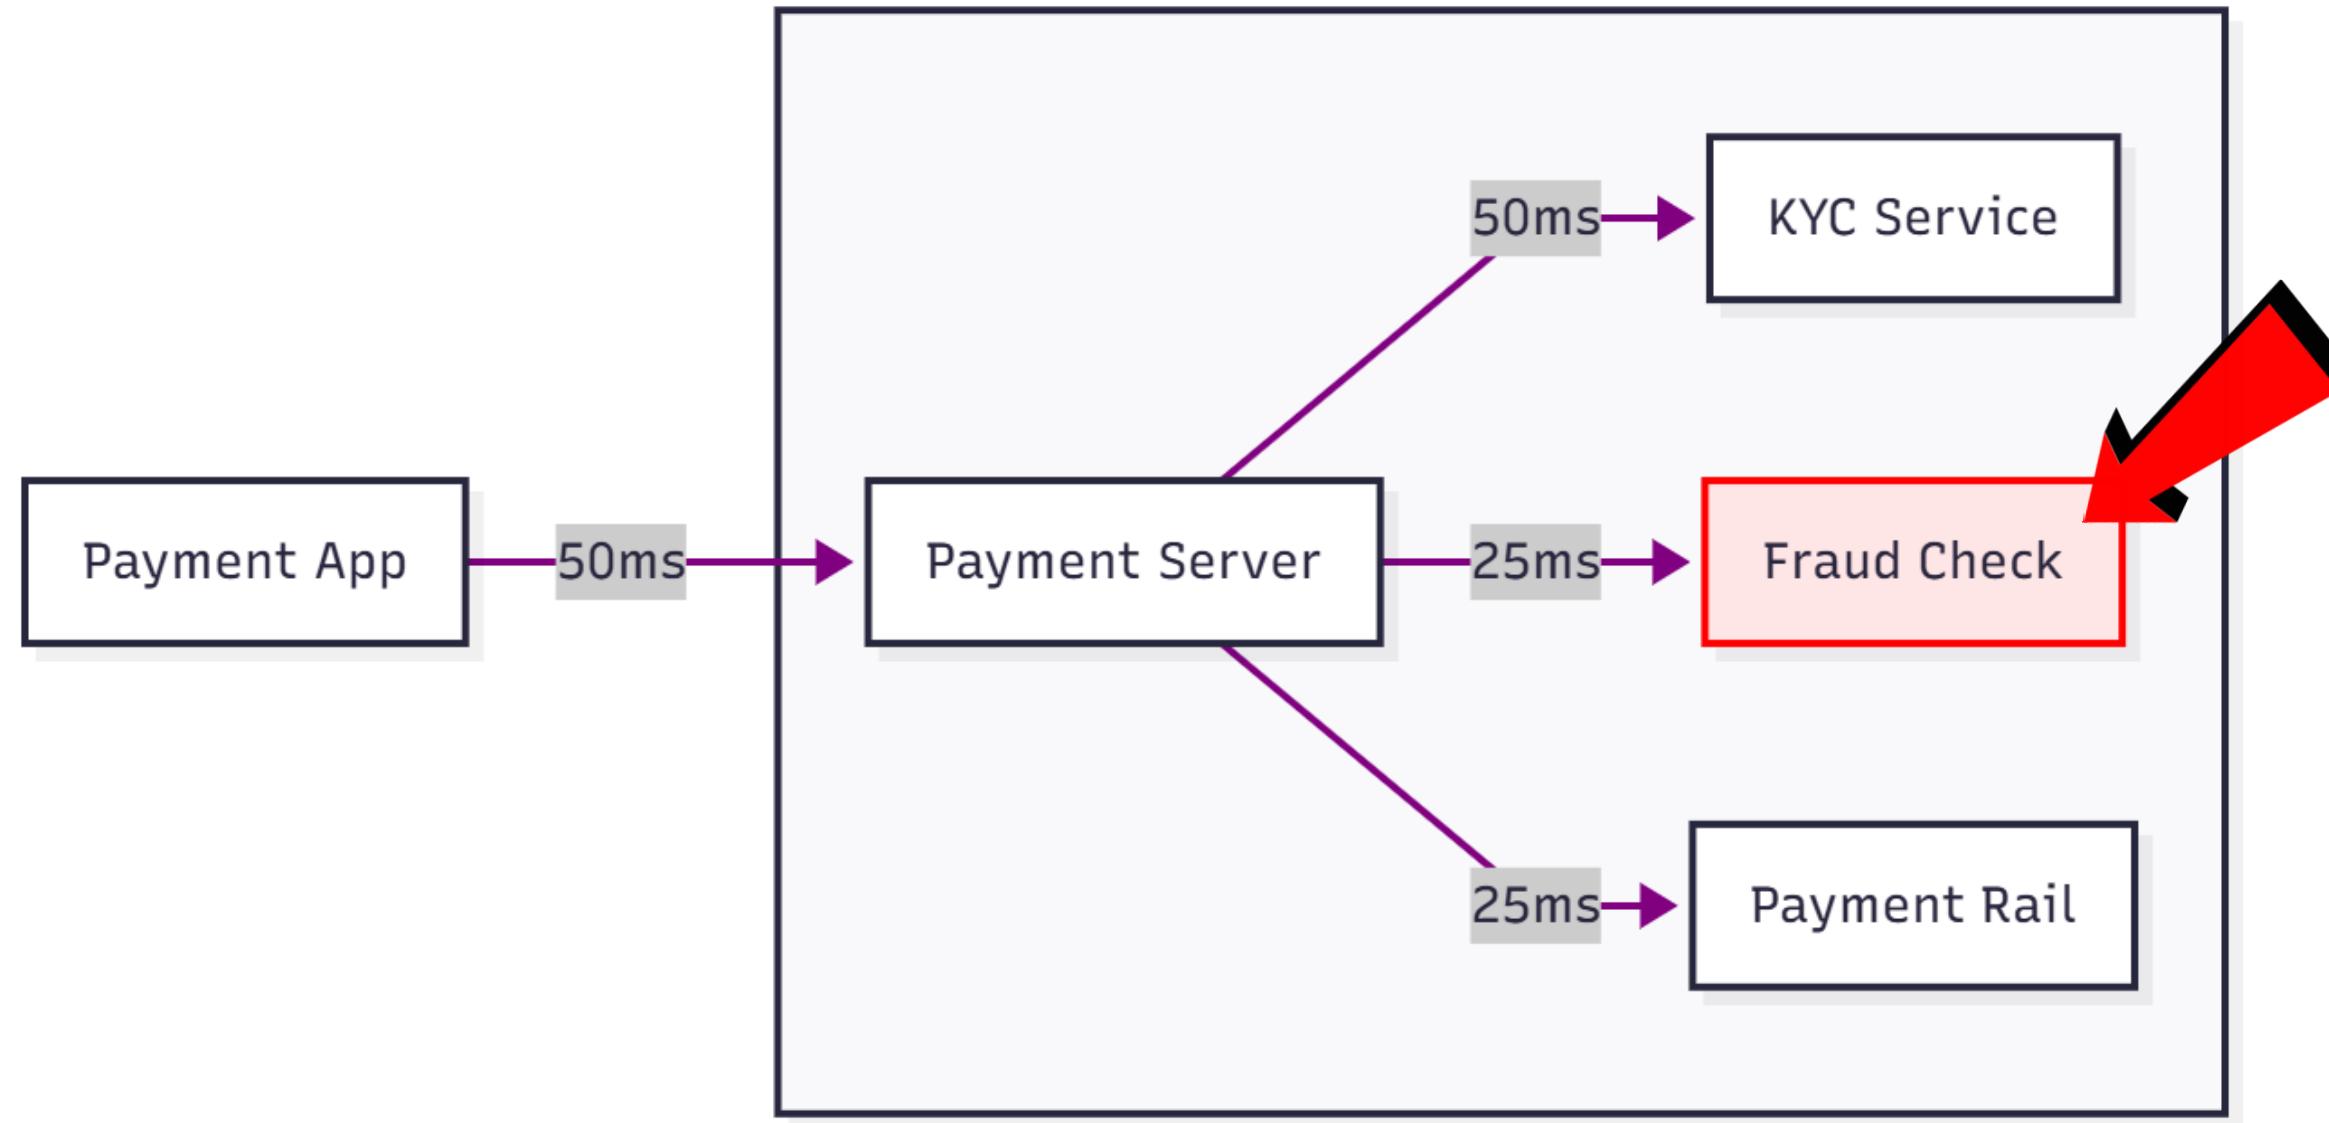
**Remember:** Random splits give you optimistic metrics that collapse in production. Time-based validation predicts real performance.

## Production-Style Rolling Window Visualization



# **4. Latency-Aware Modeling**

# Latency



# Latency Budgets: Rules vs ML

<1ms

## Rule-Based Systems

Blazingly fast, deterministic, explainable. Perfect for known patterns and hard constraints like blocklists or velocity checks.

10-200ms

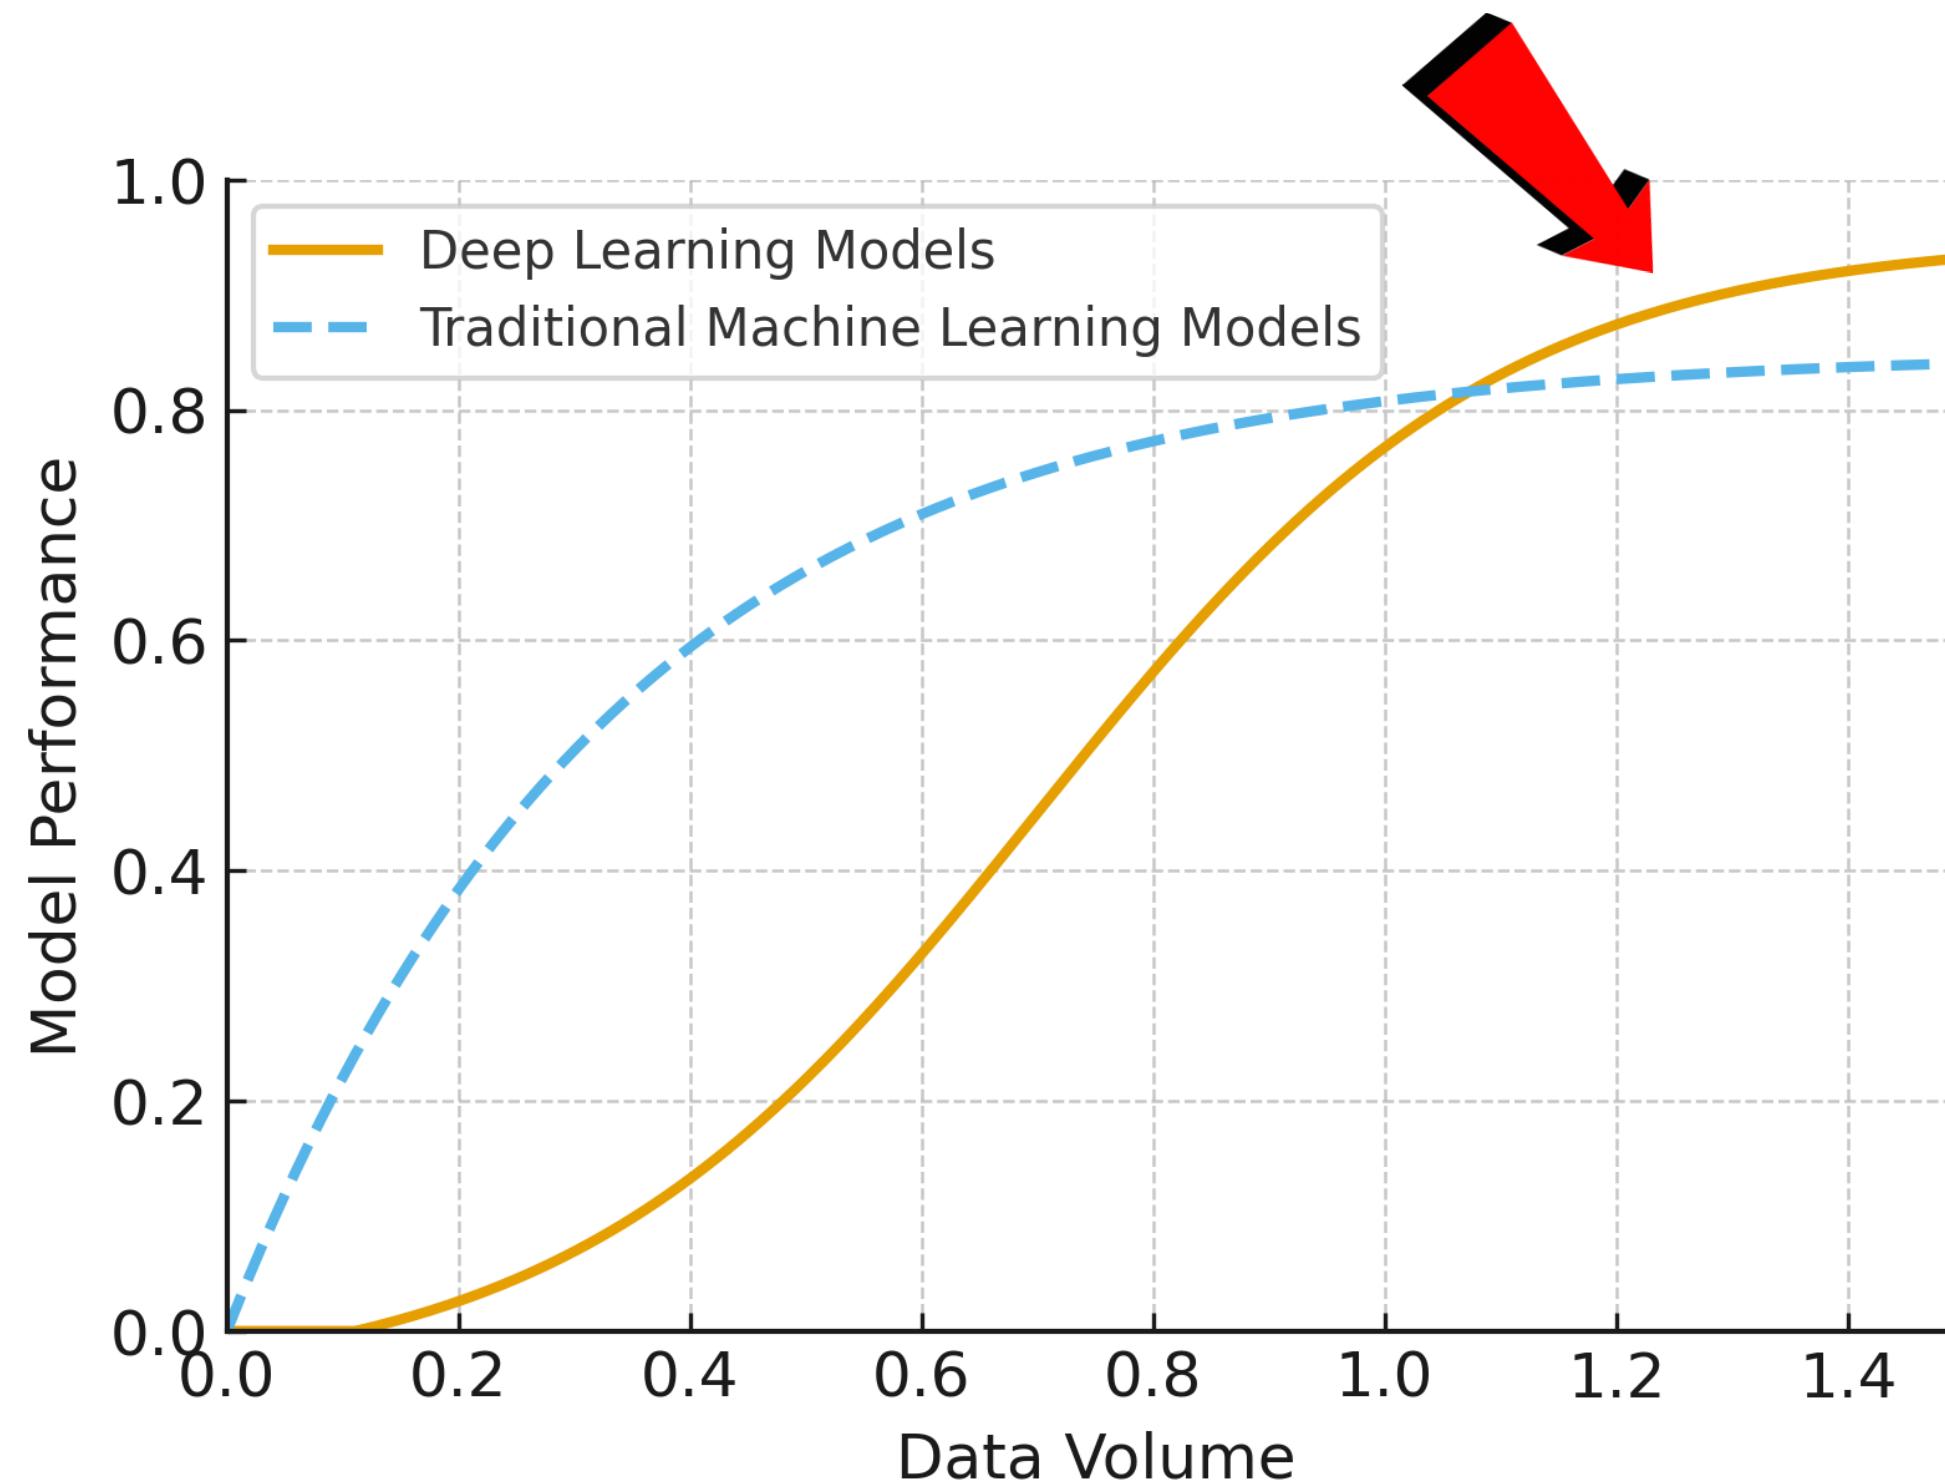
## ML Model Inference

Gradient-boosted trees typically run in tens of milliseconds. Neural networks can reach 100-200ms depending on complexity and infrastructure.

~

## Hybrid Architecture

Blend rules for instant failsafes and ML for incremental lift. Rules catch obvious fraud at <1ms, models score the nuanced middle at acceptable latency.



# Start Strong With XGBoost

## Why XGBoost Shines on Tabular Data

- Handles heterogeneous features naturally without complex preprocessing
- Fast inference fits realtime latency budgets with sub-100ms predictions
- Interpretable feature importances and tree structures support debugging and stakeholder trust
- Robust to missing values and outliers common in production fraud data

**Pro tip:** Start here before exploring deep learning. XGBoost delivers 80% of the value with 20% of the complexity.

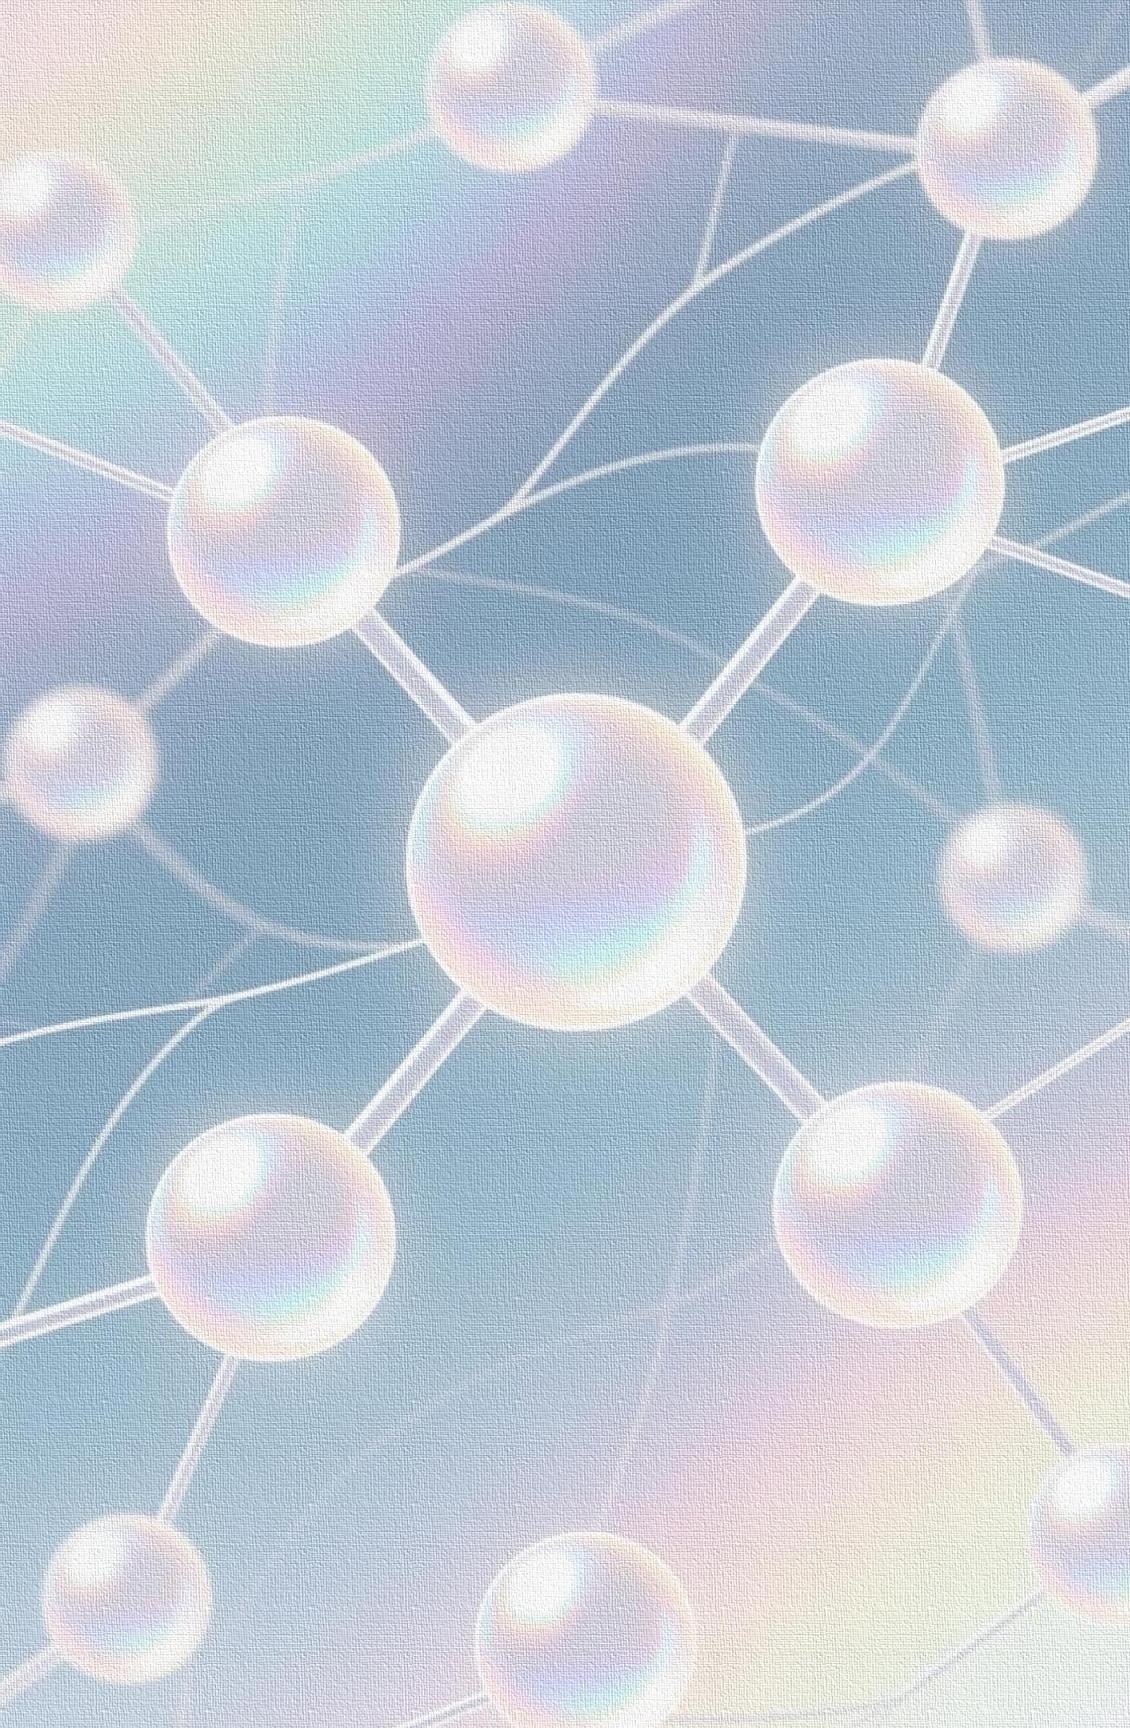


## XGBoost Python Package

This page contains links to all the python related documents on python package. To install the package, checkout [Installation Guide](#).

### Contents

- [Python Package Introduction](#)
  - [Install XGBoost](#)
  - [Data Interface](#)
  - [Setting Parameters](#)
  - [Training](#)
  - [Early Stopping](#)
  - [Prediction](#)
  - [Plotting](#)
  - [Scikit-Learn interface](#)
- [Using the Scikit-Learn Estimator Interface](#)
  - [Overview](#)
  - [Early Stopping](#)
  - [Obtaining the native booster object](#)
  - [Prediction](#)
  - [Number of parallel threads](#)
- [Python API Reference](#)
  - [Global Configuration](#)
    - [config\\_context\(\)](#)
    - [set\\_config\(\)](#)
    - [get\\_config\(\)](#)
    - [build\\_info\(\)](#)
  - [Core Data Structure](#)
    - [DMatrix](#)



# When to Add Deep Learning

Deep learning shines when other models hit a ceiling. If the feature space is sparse and high-cardinality, then tabular architectures like TabNet or FT-Transformer can capture complex patterns that gradient boosting misses.

## **Sparse, high-cardinality features benefit**

Embeddings compress categorical explosions into dense representations, enabling models to learn nuanced entity relationships.

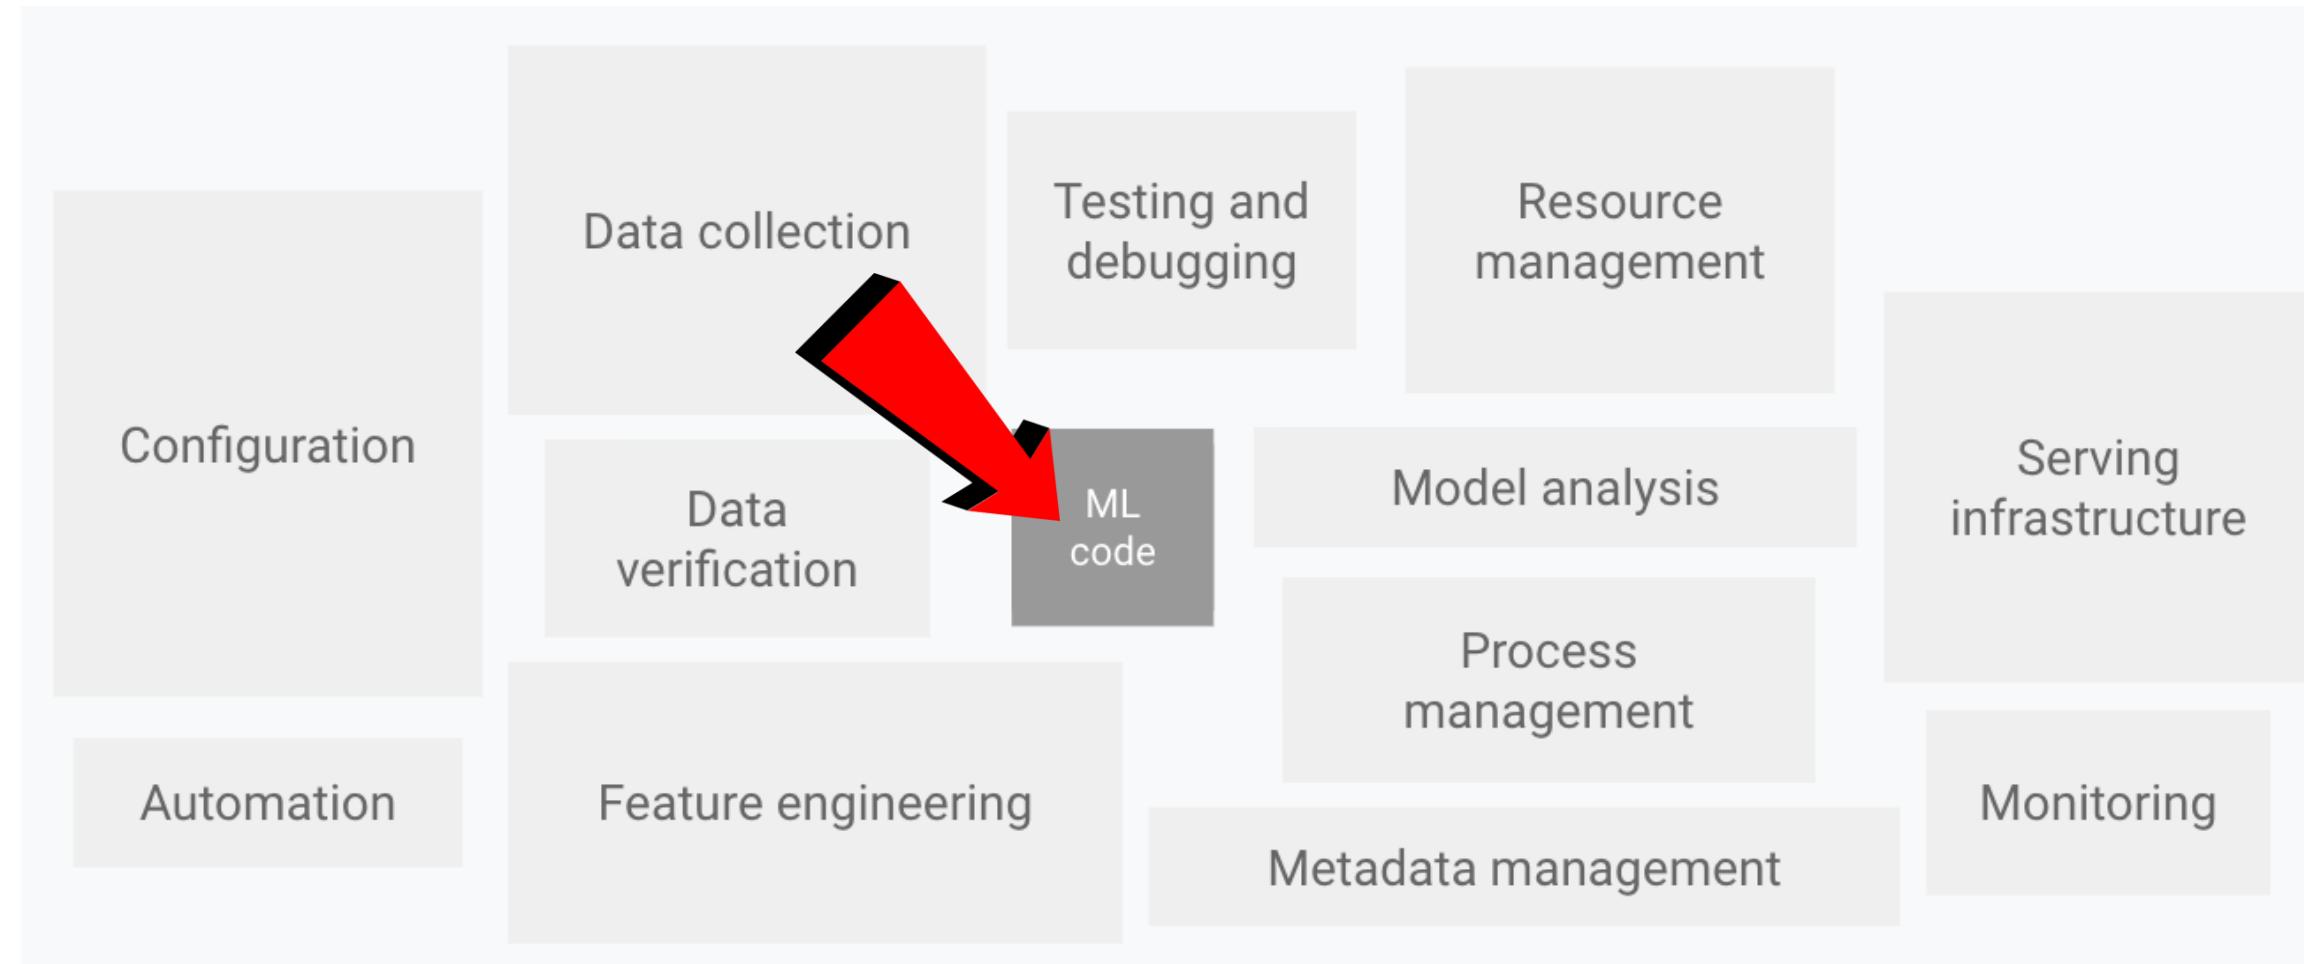
## **Learn interactions automatically**

Deep architectures discover feature crosses without manual engineering, reducing technical debt and iteration cycles.

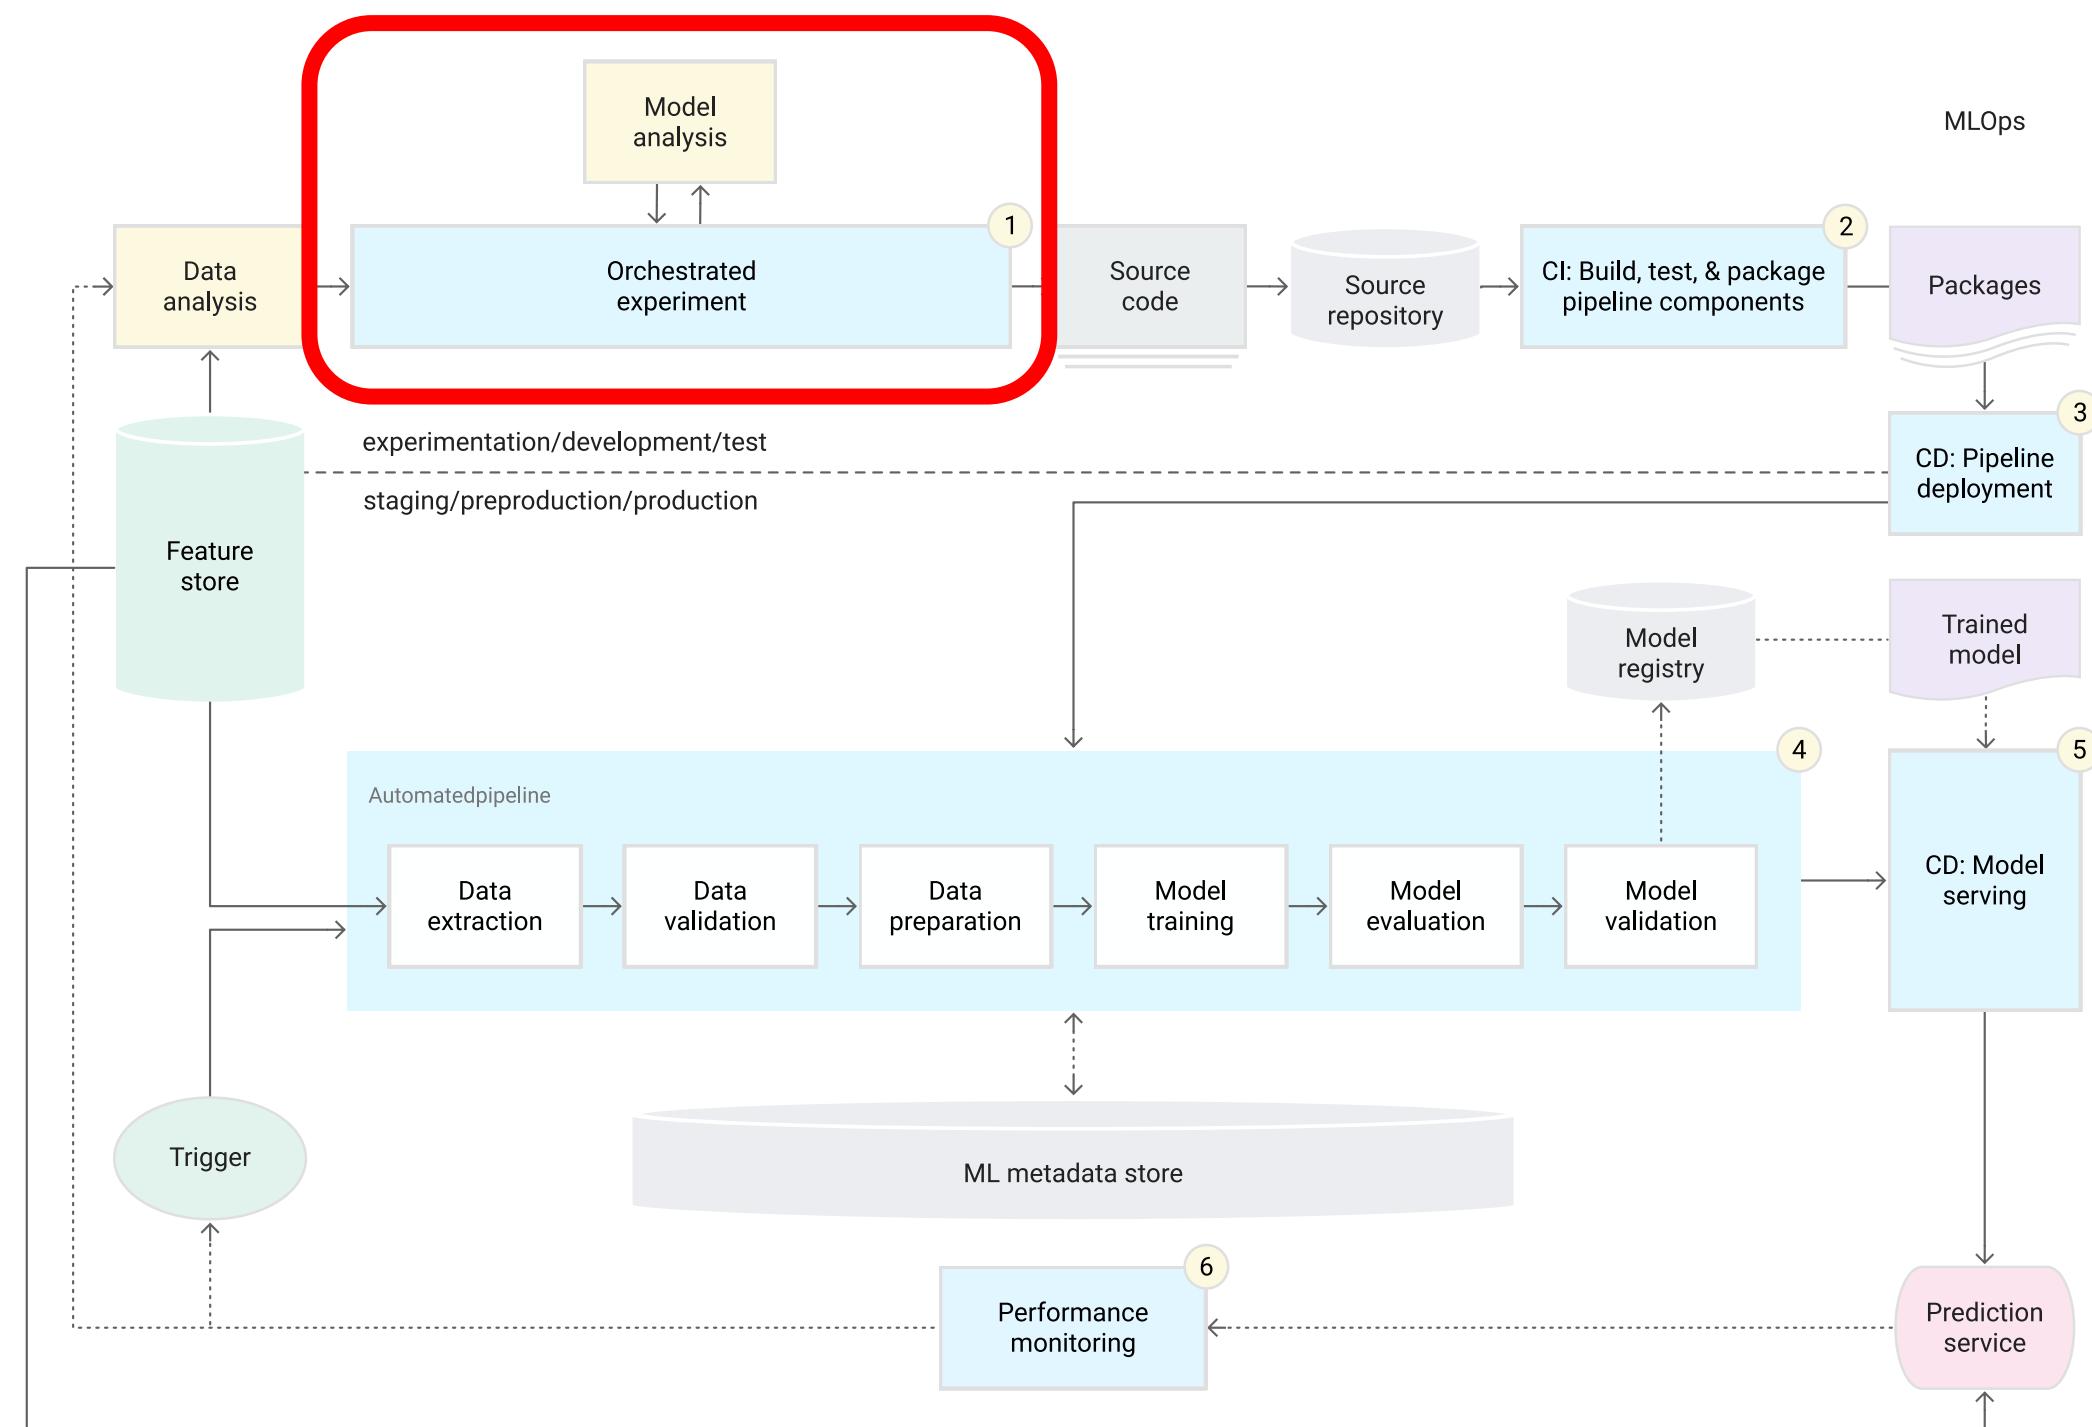
## **Validate gains against latency budget**

Measure P99 inference time in production. If your scoring SLA is 50ms, confirm DL delivers ROI before scaling.

# **5. Shipping & Operations**

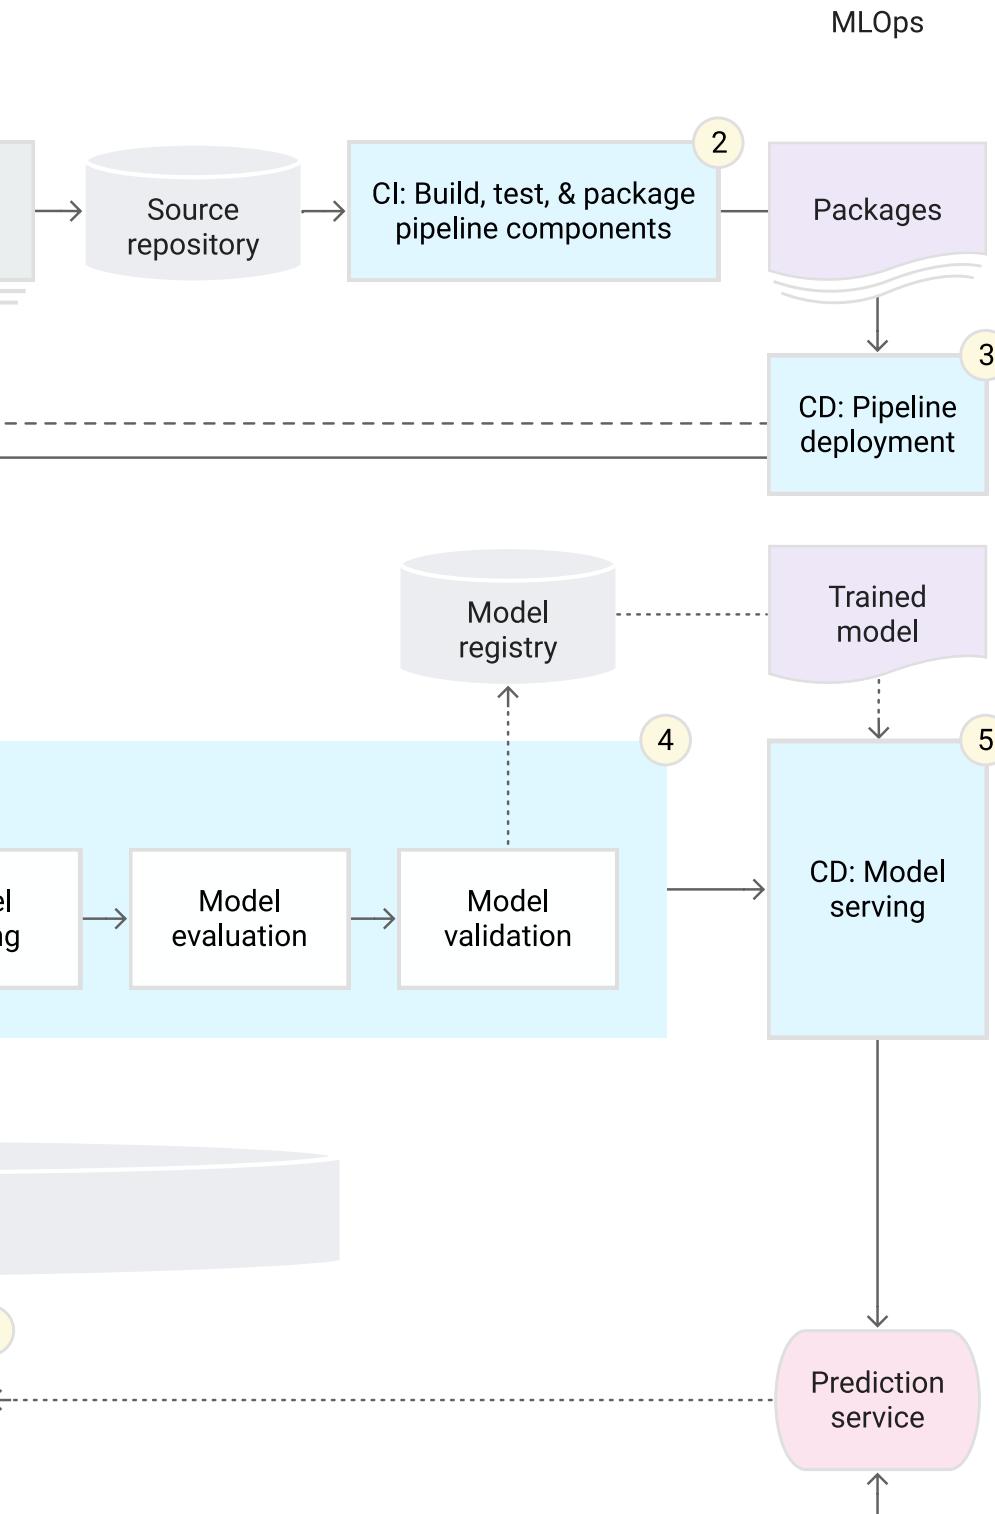


Source: <https://docs.cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>



Source: <https://docs.cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning>

# From Notebook to Service



Production fraud detection demands more than a trained model. The gap between research code and a resilient service is wide.

## 1

### Small, testable scoring core

Isolate inference logic into a pure function: features in, prediction out. Unit test edge cases (e.g., null fields, out-of-vocabulary tokens, schema drift).

## 2

### Package model plus preprocessors

Bundle your serialized model with feature transformers, scalers, and encoders. Version everything together to prevent train-serve skew.

## 3

### Serve behind a FastAPI endpoint

Wrap your scoring function in a lightweight REST API. FastAPI provides async support, automatic docs, and validation, which is critical for uptime and debugging.

# Serving Patterns That Scale

No single serving architecture fits every fraud use case. Match your pattern to transaction velocity, latency requirements, and operational complexity. Hybrid approaches often win in production.



## Batch scoring for nightly sweeps

Score historical transactions offline. Perfect for investigating past fraud rings or recomputing risk for dormant accounts (no real-time pressure).



## Online scoring for checkout flows

Synchronous prediction at transaction time. Requires sub-100ms P99 latency. Cache aggressively, use feature stores, and pre-warm models.



## TorchServe or TFServing when needed

Leverage framework-native serving for GPU inference, model versioning, and A/B testing.

# Patterns That Survive Reality

Fraud detection in production is a systems problem, not just a modeling problem. The best ML model fails if it can't ship, scale, or adapt.

## Decide for cost, not vanity scores

Optimize for business impact (false positive costs, analyst capacity, customer friction). Precision-recall tradeoffs are product decisions.

## Validate in time, always

Temporal cross-validation prevents overfitting to old fraud patterns. Adversaries evolve and your evaluation strategy must reflect that reality.

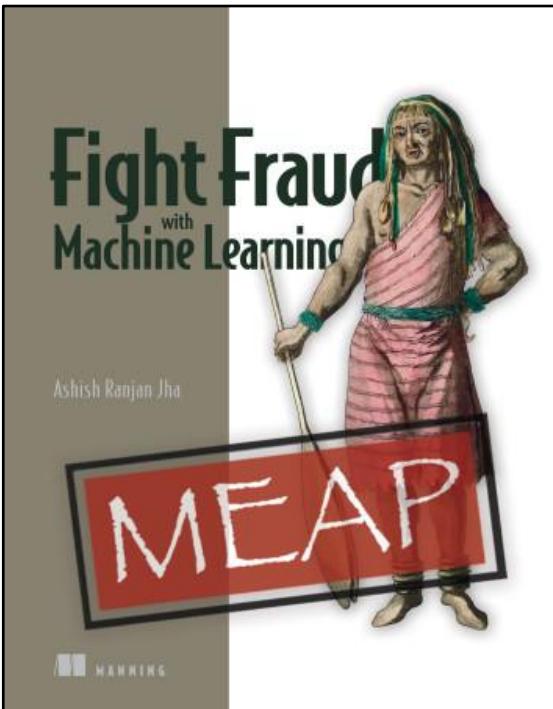
## Ship with guardrails, monitor relentlessly

Use a hybrid approach, add complexity to the model only if necessary, and stick to MLOps good practices. Observability is not optional.



# Recommended Resources

1. "Fight Fraud with Machine Learning", MEAP book by Manning
2. <https://github.com/safe-graph/graph-fraud-detection-papers/>



## Awesome Graph Fraud Detection

A curated list of graph-based fraud, anomaly, and outlier detection papers & resources

Table of Content

- Deep Learning Papers: [2025](#) | [2024](#) | [2023](#) | [2022](#) | [2021](#) | [2020](#) | [Before 2020](#)
- [Non-Deep-Learning Papers since 2014](#)
- [Toolbox](#)
- [Dataset](#)
- [Survey Paper](#)
- [Other Resource](#)

# Q&A