




Bodo DataFrames: a fast and scalable HPC-based drop-in replacement for Pandas

PyData Global, December 11, 2025

Scott Routledge
Software Engineer
Bodo.ai



About Me


- Software Engineer at Bodo.ai
- B.S. in CS and ML @ Carnegie Mellon
- Worked on Python Compiler & BodoSQL
- Currently working on Bodo DataFrames





High Performance Data Processing in Python

- Open source:
 - `pip install bodo`
 - `conda install bodo -c conda-forge`
 - `import bodo.pandas as pd`
- Parallel computing with MPI
- JIT compilation & query optimization
- SQL engine (`pip install bodosql`)



Agenda

- ✓ The Challenge: Simple, Scalable Data Infrastructure
- ✓ Bodo DataFrames Overview and Demo
- ✓ JIT Compiler for Custom Code
- ✓ Integration with Iceberg Table Format

The Challenge: Simple Scalable Data Infra

Challenges:

- ✗ Frequent code rewrites and performance tuning
- ✗ Complexity in managing large cloud datasets
- ✗ Fragmentation of tools for different use cases

Requirements:

- ✓ Fast and scalable for data size without code changes
- ✓ Easy to use and understand
- ✓ Flexibility: ETL, analytics, preprocessing, AI workloads (inference, training)

Bodo DataFrames

- ✓ Database-grade query optimizer & JIT compiler
- ✓ MPI backend with streaming execution
- ✓ Spill to disk to prevent Out of Memory (OOM) Errors
- ✓ Fallback to Pandas for unsupported operations

```
import bodo.pandas as pd

df = pd.read_parquet(path)
df["month"] = df.F.dt.month
df["year"] = df.F.dt.year
df = df[df.B == "gamma"]

df.to_parquet("result.pq")
```

Demo Notebook

NYC Taxi ETL




Bodo Scales Pandas Efficiently

Could use Spark, Dask, etc.


- Implementation complexity
- Task-based model creates high overheads

Or... Use Bodo!

- Drop-in Pandas replacement
- HPC execution scales from laptop to large cluster
- JIT compilation & query optimization



DataFrames Workflow



DataFrames Workflow

User process

```
import bodo.pandas as pd\n\ndf = pd.read_parquet(path)\n\ndf[\"month\"] = df.F.dt.month\n\ndf[\"year\"] = df.F.dt.year\n\ndf = df[df.B == \"gamma\"]\n\ndf.to_parquet(\"result.pq\")
```

Generated Plan

PROJECTION
Expressions:
...
"year" (#[1..3])

PROJECTION
Expressions:
...
"month" (#[0..3])

BODO_READ_PARQUET(A, B...
~10 Rows

DataFrames Workflow

User process

```
import bodo.pandas as pd\n\ndf = pd.read_parquet(path)\n\ndf[\"month\"] = df.F.dt.month\n\ndf[\"year\"] = df.F.dt.year\n\ndf = df[df.B == \"gamma\"]\n\ndf.to_parquet(\"result.pq\")
```

Generated Plan


FILTER
Expressions:
(#[df.A] = 'gamma')

PROJECTION
Expressions:
...

PROJECTION
Expressions:
...


BODO_READ_PARQUET(A, B...

DataFrames Workflow




Plan Optimization: Filter Pushdown

```
import bodo.pandas as pd  
  
df = pd.read_parquet(path)  
  
df["month"] = df.F.dt.month  
  
df["year"] = df.F.dt.year  
  
df = df[df.B == "gamma"]
```




Vectorized Execution Example

```
import bodo.pandas as pd\n\ndf = pd.read_parquet("raw_data.pq")\n\ndf["ts"] = df.ts.dt.ceil("100ms")\n\ndf.groupby("sensor_id", as_index=False\n).agg({"value": "mean", "ts": "min"})\n\ndf.to_parquet("data_agg.pq")
```




Vectorized Execution Example



- Logical plan is converted to a series of pipelines
- Push based execution model: data flows from source node(s) to a sink

Async Messaging in Pipelines



- Workers process pipelines independently
- Async MPI communication can overlap with compute
- Workers sync only at pipeline boundaries

DataFrame Libraries Comparison

	Ease of Use	Query Optimizer	Execution
Bodo DataFrames	Full Pandas compatibility with fallback	Database-grade optimizer	HPC/MPI and database-based backend
Pandas on Spark	Pandas compatible; JVM complexity	Spark Catalyst (database-grade)	Task scheduling on JVM
Dask	Mostly Pandas compatible; explicit “.compute()” and code changes	Basic compute graph optimization	Task scheduling in Python
Modin/Ray	Pandas-compatible with fallbacks	No optimizer	Slow task scheduling (Ray-based)
Polars	SQL-like API, not Pandas-compatible	Simple heuristic rules	Single-node threaded task scheduling
Daft	SQL-like API, not Pandas-compatible	Simple heuristic rules	Ray-based task scheduling (non-vectorized)

JIT Compilation for Custom Code

- Bodo DataFrames tries JIT compilation for ser.map / df.apply calls
- Falls back to Python if compilation fails

```
def simulate(row):  
    balance = row.start_balance  
    for _ in range(row.steps):  
        balance += random.randint(-1, 1) * row.drift  
    return balance  
  
df["final_balance"] = df.apply(simulate, axis=1)
```

Bodo JIT Compiler

```
@bodo.jit
def example():
    table = pd.read_parquet('data_parquet')
    data = table[table['A'].str.contains('ABC*', regex=True)]
    stats = data['B'].describe()
    print(stats)
```


- ✓ Compiles Pandas/Numpy to optimized binary with MPI automatically
- ✓ Automatically applied to UDFs of Bodo DataFrames
- ✓ Alternative to Bodo DataFrames for higher performance for custom code
- ✗ Limitations: Compilation delays, unsupported ops, no streaming/spilling support

Benchmarks

NYC Taxi (~1B rows)


SINGLE NODE PERFORMANCE COMPARISON

(1 x r6i.16xlarge)



CLUSTER PERFORMANCE COMPARISON

(4 x r6i.16xlarge)



Bodo Iceberg Support

- Supports `read_iceberg` and `to_iceberg` (new Pandas 3.0 APIs)
- Easy set up: single line of Python code to write or read
- Query past versions of datasets via time travel
- Leverages iceberg's metadata for better query planning
- Filter pushdown becomes even faster with Iceberg!


```
import bodo.pandas as pd

df = pd.read_iceberg(
    table_name, location=table_bucket,
    snapshot_id=prev_snapshot
)
filtered_df = df[df.ts > CUTOFF_DATE]
```

Summary


Bodo provides high-performance data processing in native Python

- True parallel computing with powerful MPI backend
- Pandas drop-in replacement
- JIT compiler & query optimizer



Thank you!

- Questions?
- Try Bodo Yourself 
 - pip install bodo
 - conda install bodo -c conda-forge
 - Join our [Slack community](#)
 - Star us on [Github](#)



Supplemental Slides (if time allows)

Plan Optimization: Join Reordering


- Joins can be applied in different orders and produce equivalent results
- Join order effects intermediate result sizes
- Bad join order can lead to >100x slower execution times!

```
df1 = returns.merge(online_orders, on="order_id")
df2 = locations.merge(df1, on="customer")
```

```
df1 = locations.merge(returns, on="customer")
df2 = online_orders.merge(df1, on="order_id")
```


Plan Optimization: Join Reordering

- ✖ Bad join ordering: intermediate result size explodes to much larger than inputs
- ✖ Difficult to predict: simple heuristics like “join the smaller tables first” don’t work



Plan Optimization: Join Reordering


- ✓ Good join ordering: reduce intermediate result size by 100x!
- ✓ Planner uses dynamic programming to explore different orderings and picks the best one based on cost estimates



Plan Optimization: Join Filters

```
orders = pd.read_parquet("orders.pq")
cust = pd.read_parquet("customers.pq")

orders.merge(cust, on="CUSTOMER_ID")
```



Plan Optimization: Join Filters

- Use runtime information from build table to filter probe table before the join
- Join filters can be pushed down to reduce intermediate results sizes
- Recently added to Bodo DataFrames

