


When the Meter Maxes Out

Chernobyl Disaster
Lessons for ML
Systems in
Production

PyData Global 2025

Idan Richman Goshen
Data Science
Team Lead @ **Lusha** 





3.6 ROENTGEN NOT GREAT, NOT TERRIBLE

The truth was thousands of times higher!

People trusted the instrument because it gave
a number. any number.

**Instrumentation is not ground truth
it is a model of the truth.**

About Me

Idan Richman Goshen

- Economist turned Data Scientist (~8 years ago)
- Coding since childhood (drawing polygons in assembly 🤖)
- Board member, Eretz-Ir NGO, making cities community-oriented
- Experience across both startups and large enterprises



Timeline of the Disaster



Before the Explosion: Critical Failures

April 25, 1986

Operational Pressure



Safety Override



Feedback Loop





- The test was delayed because the electrical grid demanded full power during the day, so they pushed ahead at night with a less experienced and tired crew.



A recommender system built in a hurry.

What went wrong?

A bug in the user's table caused all user features to become null, and your model failed silently.



Core lessons

- Speak up**
- If you end up compromising, make sure all stakeholders are aware of the risks**
- Expect the worst - prepare fallbacks**
- “Don’t deploy on Friday” 😊**

Safety Override



The operators disabled the Emergency Core Cooling System, automatic shutdown logic, and several alarms “just for this test”.



Forecasting electricity demand during a severe storm.
Demand and supply patterns break, lines start tripping - all
those alerts fire continuously.

What went wrong?

The very alarms designed to catch these events were muted because everything looked abnormal already.

Even in expected special cases - like Black Friday or a storm - we can predict **only part** of the abnormal behavior.

Core lesson

Special operating modes need *their own* safety envelope:

- What “expected abnormal” looks like
- What signals remain non-negotiable
- Hard limit conditions for stop/revert





If you must bypass something,
document it, announce it, and create compensating safeguards.

Feedback Loop



- The RBMK's positive void coefficient made it unstable at low power - the very regime of the turbine test.



A NER Transformer model deployed in "Shadow Mode" with stable P99 latency at 900ms.

What went wrong?

That 900ms latency frustrates users. They hit refresh
-> one request becomes two, causing more timeouts
and more refreshes. The system enters a positive
feedback loop of retry storms

Core lessons

- Your model isn't an isolated component
- Ponder second-order interactions
- Use Canary deployments for a more realistic “Closed Loop” system (expose 1% of real users to the new model).

The Explosion: The Point of No Return

April 26, 01:23 AM



Failed Fallbacks



False Metric




The Unexpected truth

Failed Fallbacks



- AZ-5 Button: Operators pressed the emergency shutdown button to stop the instability. But under the test conditions, it acted as a detonator.



A new model is deployed, with the legacy service set as fallback.

What went wrong?

An upstream missing feature kills the new model, dropping 100% of traffic to the fallback. The legacy service - now running on a single micro-instance - is hit with a tsunami of requests, and crashes instantly.




Core lessons

- Don't plan only for the happy path**
- Account for all dependencies your system rely on**
- Be in sync with all cross-team dependencies**
- Consider simpler and maintainable fallback solutions**



A Quick Test

When you see the following, how do you react?



A timestamp feature with value of 1970-01-01



An ID field showing -1



A probability predicted as exactly 1.0 or 0.0






- The control room dosimeters maxed out at 3.6 R/h, while in reality it was around 15,000 R/h.



Fraud detection model dashboard shows a **stellar AUC around 0.95**. Meanwhile, in reality, you are losing \$100k an hour.

What went wrong?

Downstream process used model scores with a threshold of 0.8. Over time model scores drifted down. Fraudsters still got higher scores than legit users, but below the threshold.



Core lessons

- A metric is a single-dimensional model of the truth
- Track multiple metrics (e.g.: Avg. model score)
- Collect and track downstream/business metrics, even if lagged




RMBK core explosion considered “impossible”... until it wasn’t.
It is an inherent design flaw that wasn’t communicated.

- “The model is stable.” → You won’t test extreme inputs
- “The system is robust.” → You won’t simulate drift
- “It can’t go wrong.” → You won’t build fallbacks



Core lessons

- Your system can fail in more ways than you expect
- Favor business metrics and common sense over theoretical explanations
- Act professionally - Communicate your model’s limitations



Pripyat's 49,000 residents' evacuation

36 hours after the explosion



Radiation levels in Pripyat were abnormal almost immediately after the explosion. Symptoms appeared and external cities reported spikes. The authorities waited for conclusive proof before acting.



Big failures rarely begin with dashboards turning red.

By the time certainty arrives - customers complaining, KPIs collapsing - you've already expanded the blast radius



Model input distribution slowly shifts




Core lesson

Professionals don't escalate when they're certain; they escalate when something looks off

To Summarize

Smart people, doing reasonable things, in a familiar environment - until the wrong conditions align and normal rules stop applying.

Often, it's not that single small mistake, it's a collection of practices that turns a mistake into a disaster.



Thank You !

How was the talk?

